

DEEP LEARNING

Recurrent Neural Networks

EXPORT CONTROLLED - This technology or software is subject to the U.S. Export Administration Regulations (EAR), (15 C.F.R. Parts 730-774). No authorization from the U.S. Department of Commerce is required for export, re-export, in-country transfer, or access EXCEPT to country group E:1 or E:2 countries/persons per Supp.1 to Part 740 of the EAR. ECCN: EAR99


Recurrent Neural Networks

- Dates back to (Rumelhart *et al.*, 1986)
- A family of neural networks for handling **sequential data**, which involves variable length inputs or outputs
- Especially, for natural language processing (NLP)


Sequential Data

- Each data point: A sequence of vectors $x(t)$, for $1 \leq t \leq \tau$
- Batch data: many sequences with different lengths τ
- Label: can be a scalar, a vector, or even a sequence
- Example
 - Sentiment analysis
 - Machine translation

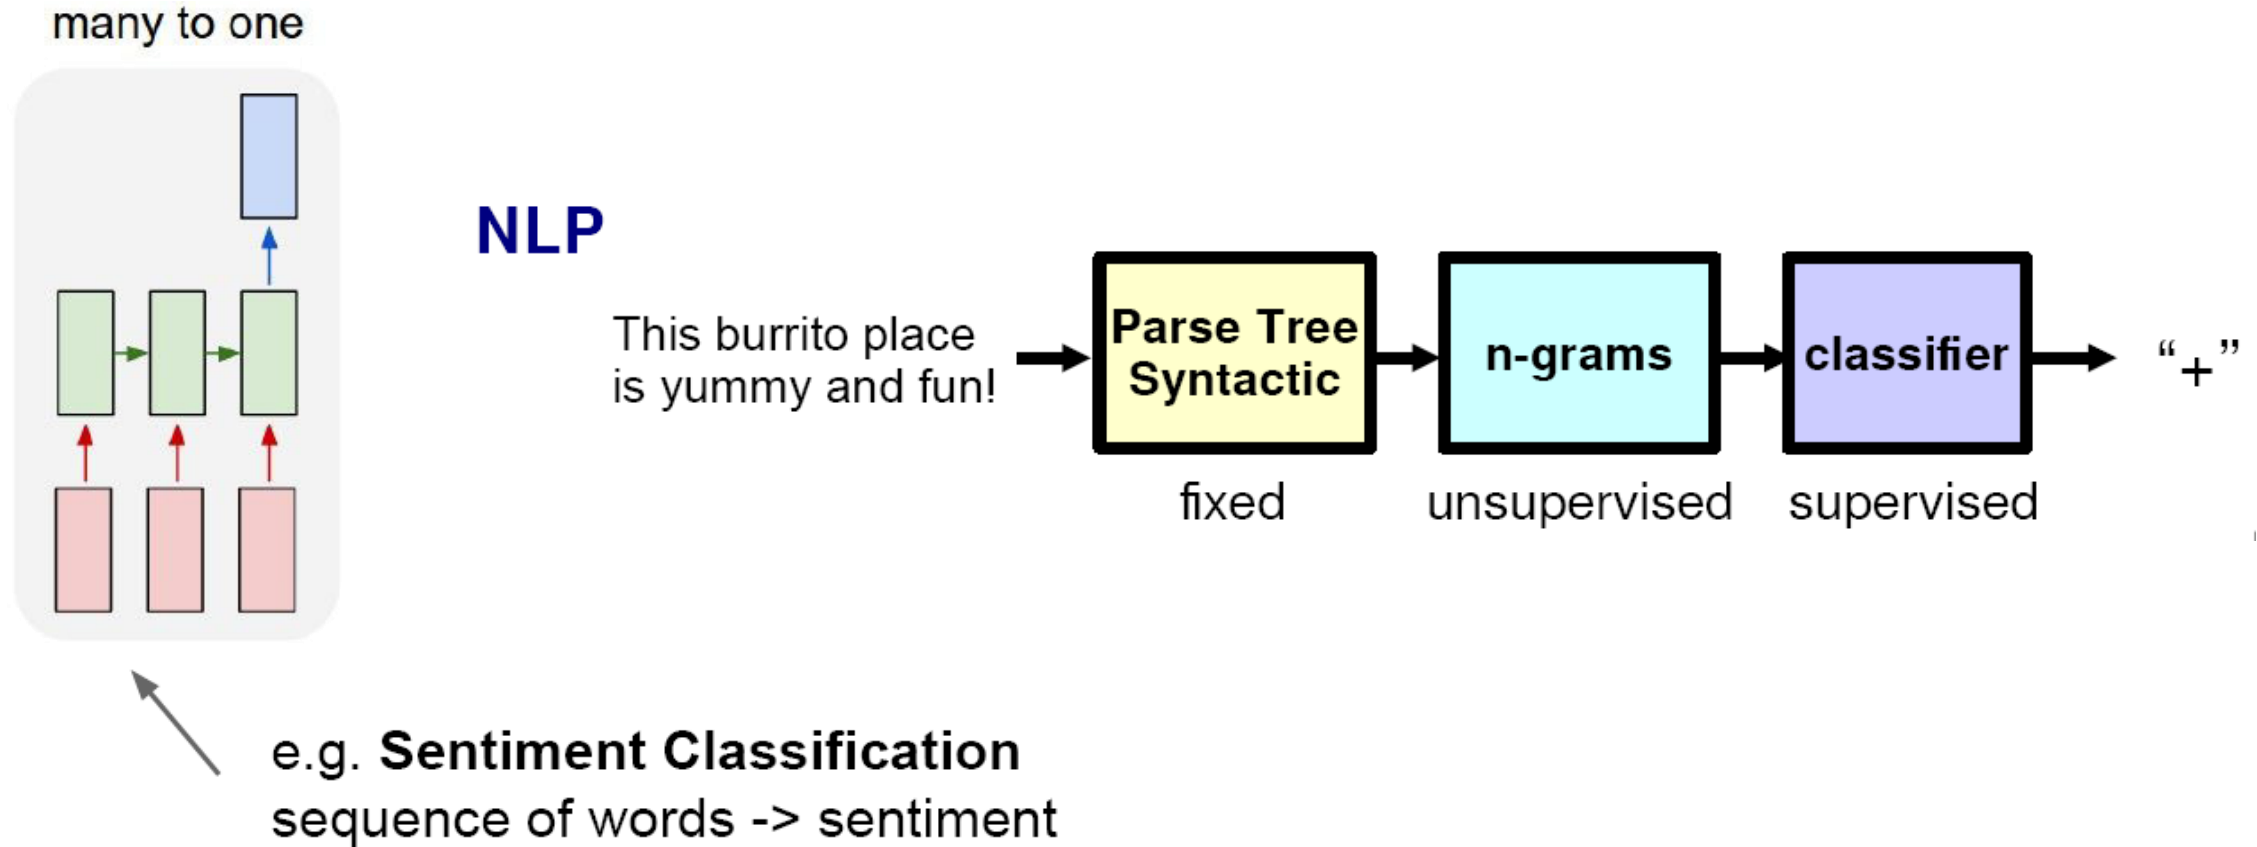
Sequential Data: Machine Translation



e.g. **Machine Translation**
seq of words -> seq of words



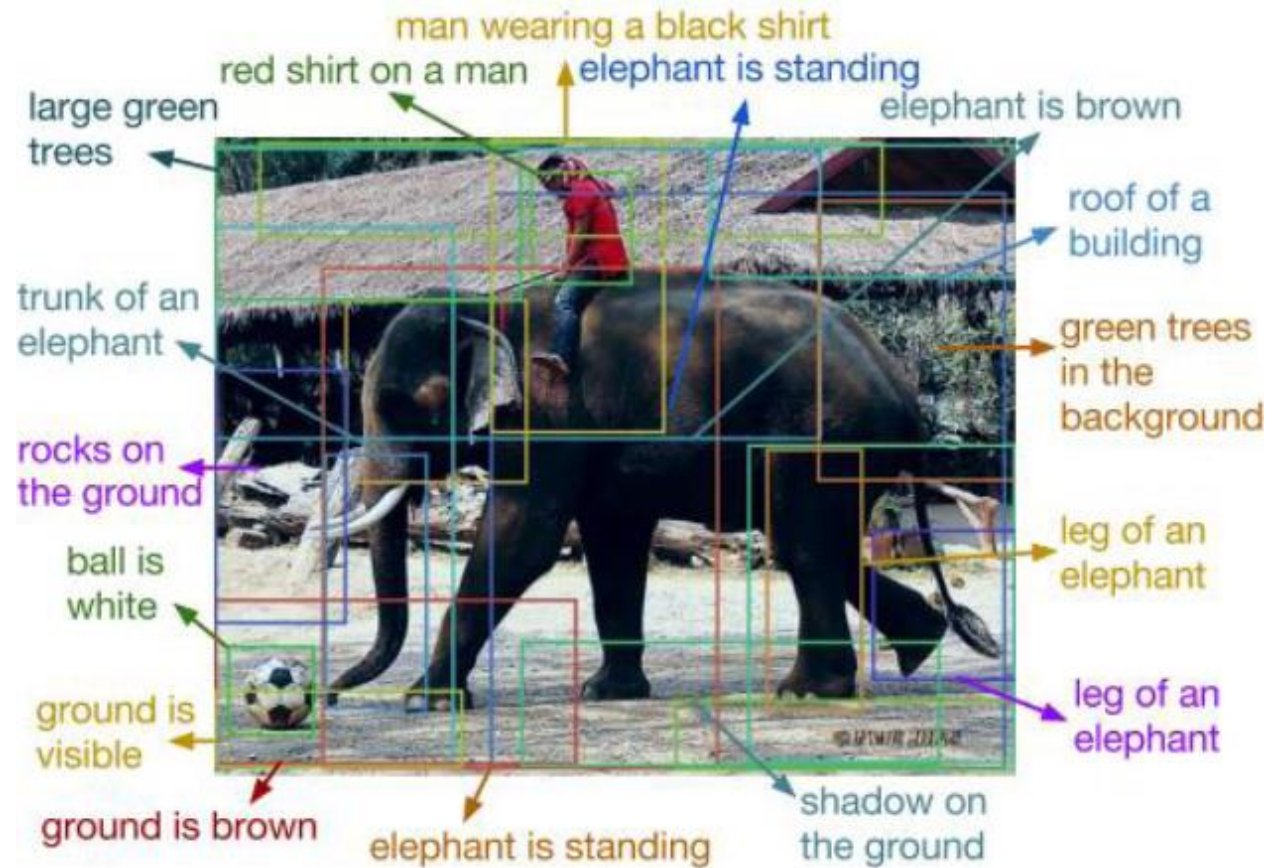
Sequential Data: Sentiment Analysis




More Complicated Sequential Data

- **Data point:** two dimensional sequences like images
- **Label:** different type of sequences like text sentences
- Example: image captioning

Image Captioning




one to many



e.g. Image Captioning
image -> sequence of words

Image source: "DenseCap: Fully Convolutional Localization Networks for Dense Captioning," by Justin Johnson, Andrej Karpathy, Li Fei-Fei


A Typical Dynamic System



$$s^{(t+1)} = f(s^{(t)}; \theta)$$

Figure from Deep Learning, Goodfellow, Bengio and Courville


A System Driven By External Data



$$s^{(t+1)} = f(s^{(t)}, x^{(t+1)}; \theta)$$

Figure from Deep Learning, Goodfellow, Bengio and Courville

Compact View (1/2)




$$s^{(t+1)} = f(s^{(t)}, x^{(t+1)}; \theta)$$

Figure from Deep Learning, Goodfellow, Bengio and Courville

Compact View (2/2)

Square: One step time delay



Key: the same f and θ for all time steps

$$s^{(t+1)} = f(s^{(t)}, x^{(t+1)}; \theta)$$


Figure from Deep Learning, Goodfellow, Bengio and Courville

Compact View

- Other forms
 - We can process a sequence of vectors x by applying a **recurrence formula** at every time step.
 - The same function and the same set of parameters are used at every time step.
 - CNNs share parameters across space; RNNs share across time.


$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state some function with parameters W old state input vector at some time step




A System Driven By External Data (1/7)

- Computational graph



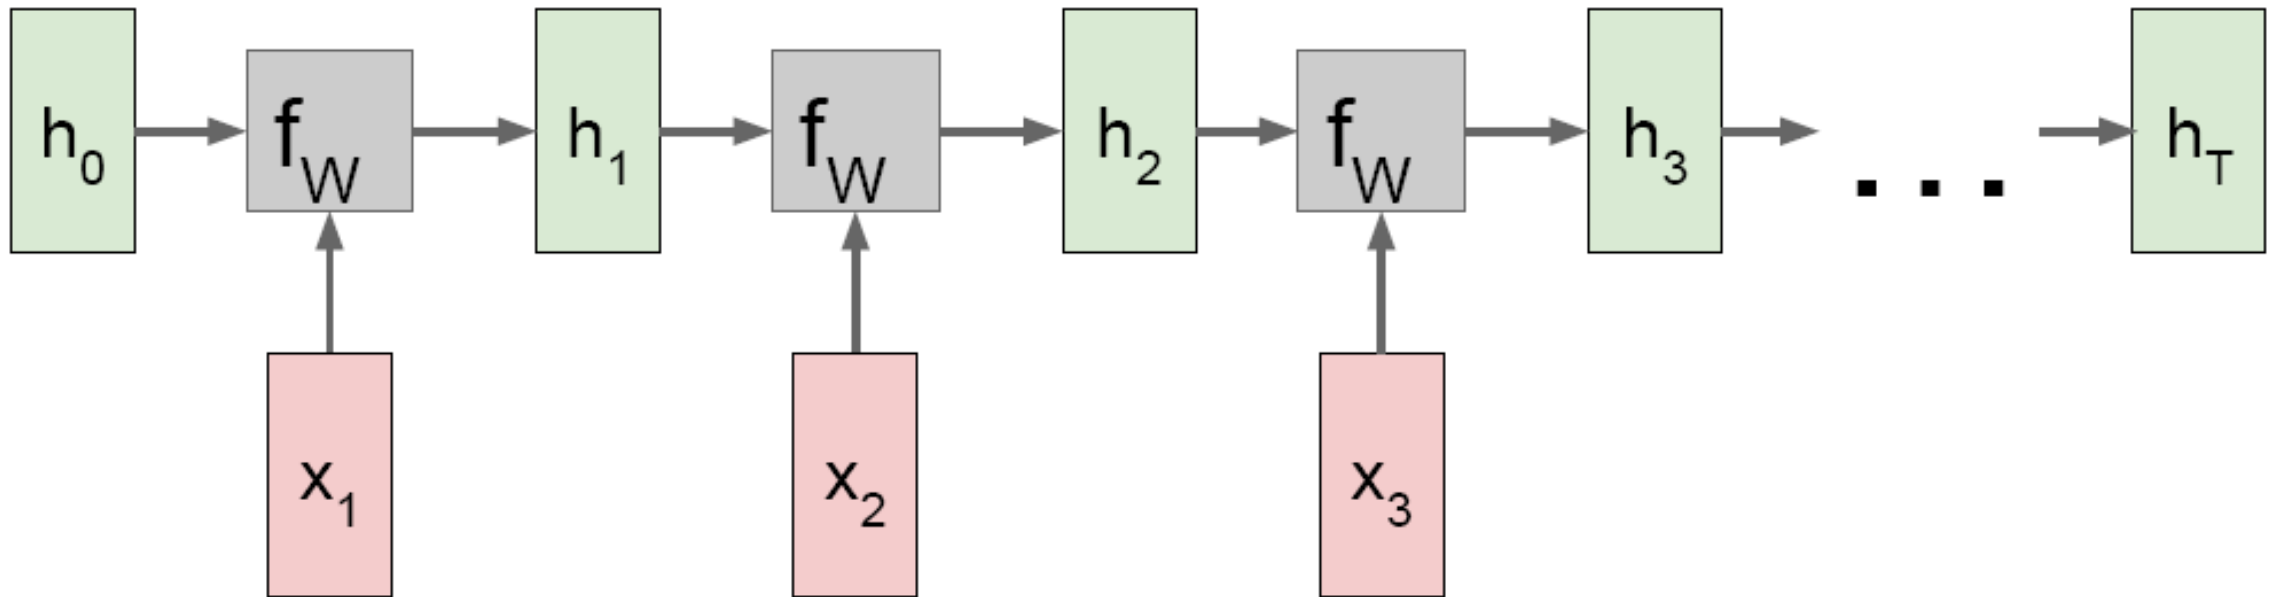
A System Driven By External Data (2/7)

- Computational graph



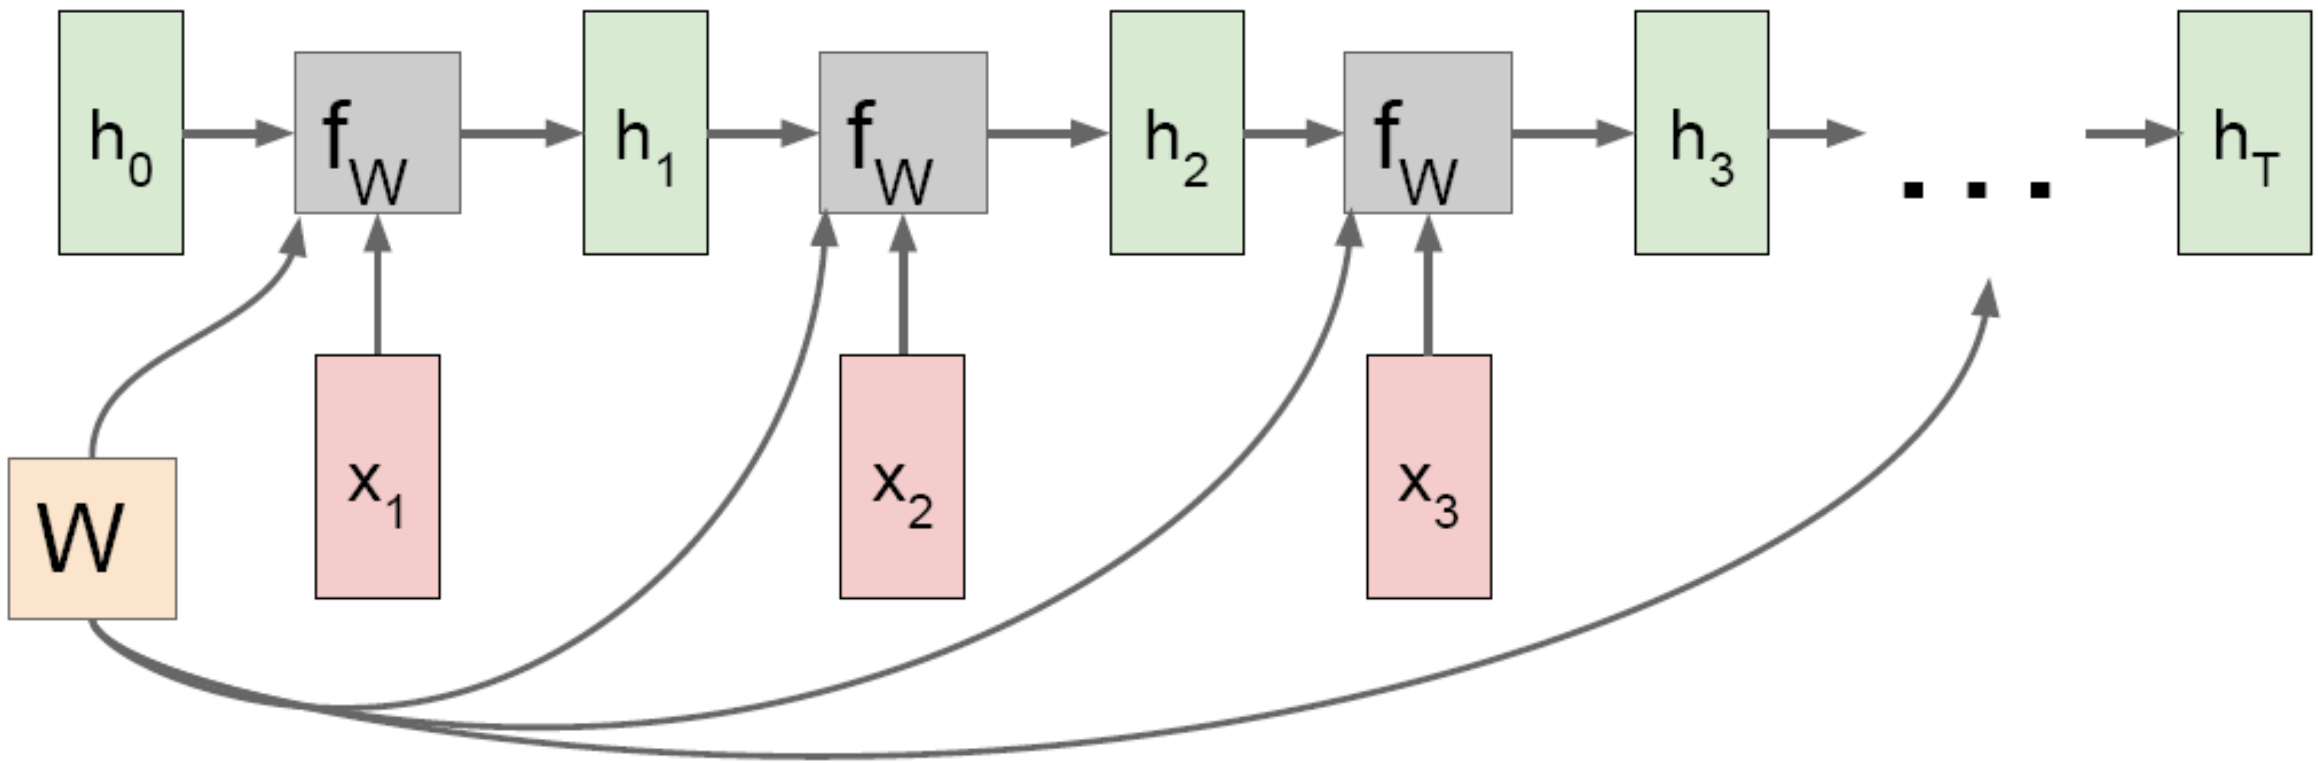
A System Driven By External Data (3/7)

- Computational graph



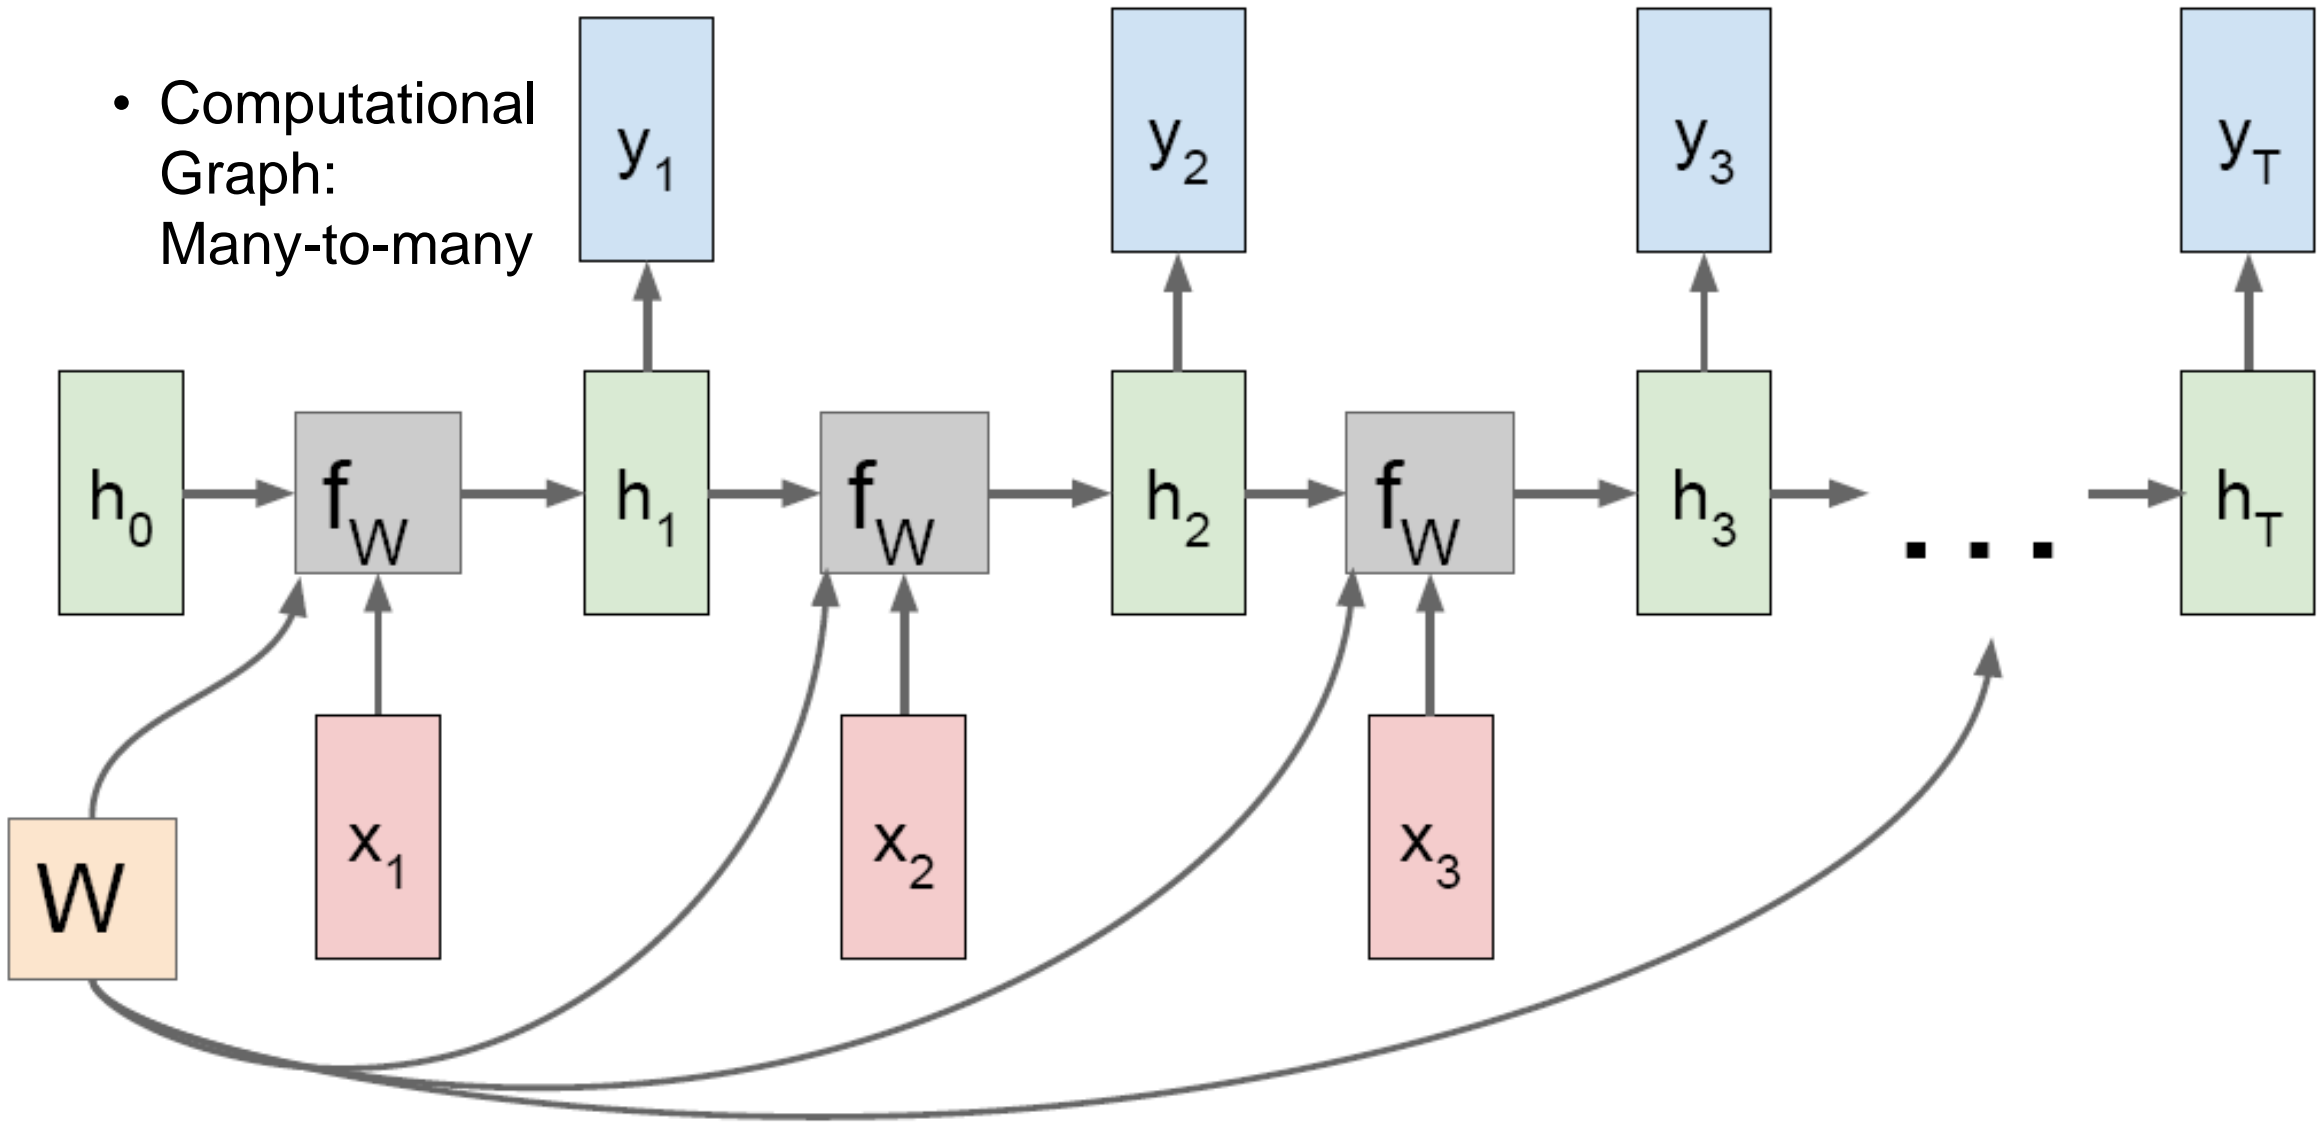
A System Driven By External Data (4/7)

- Computational graph
 - Re-use the same weight matrix at every time-step



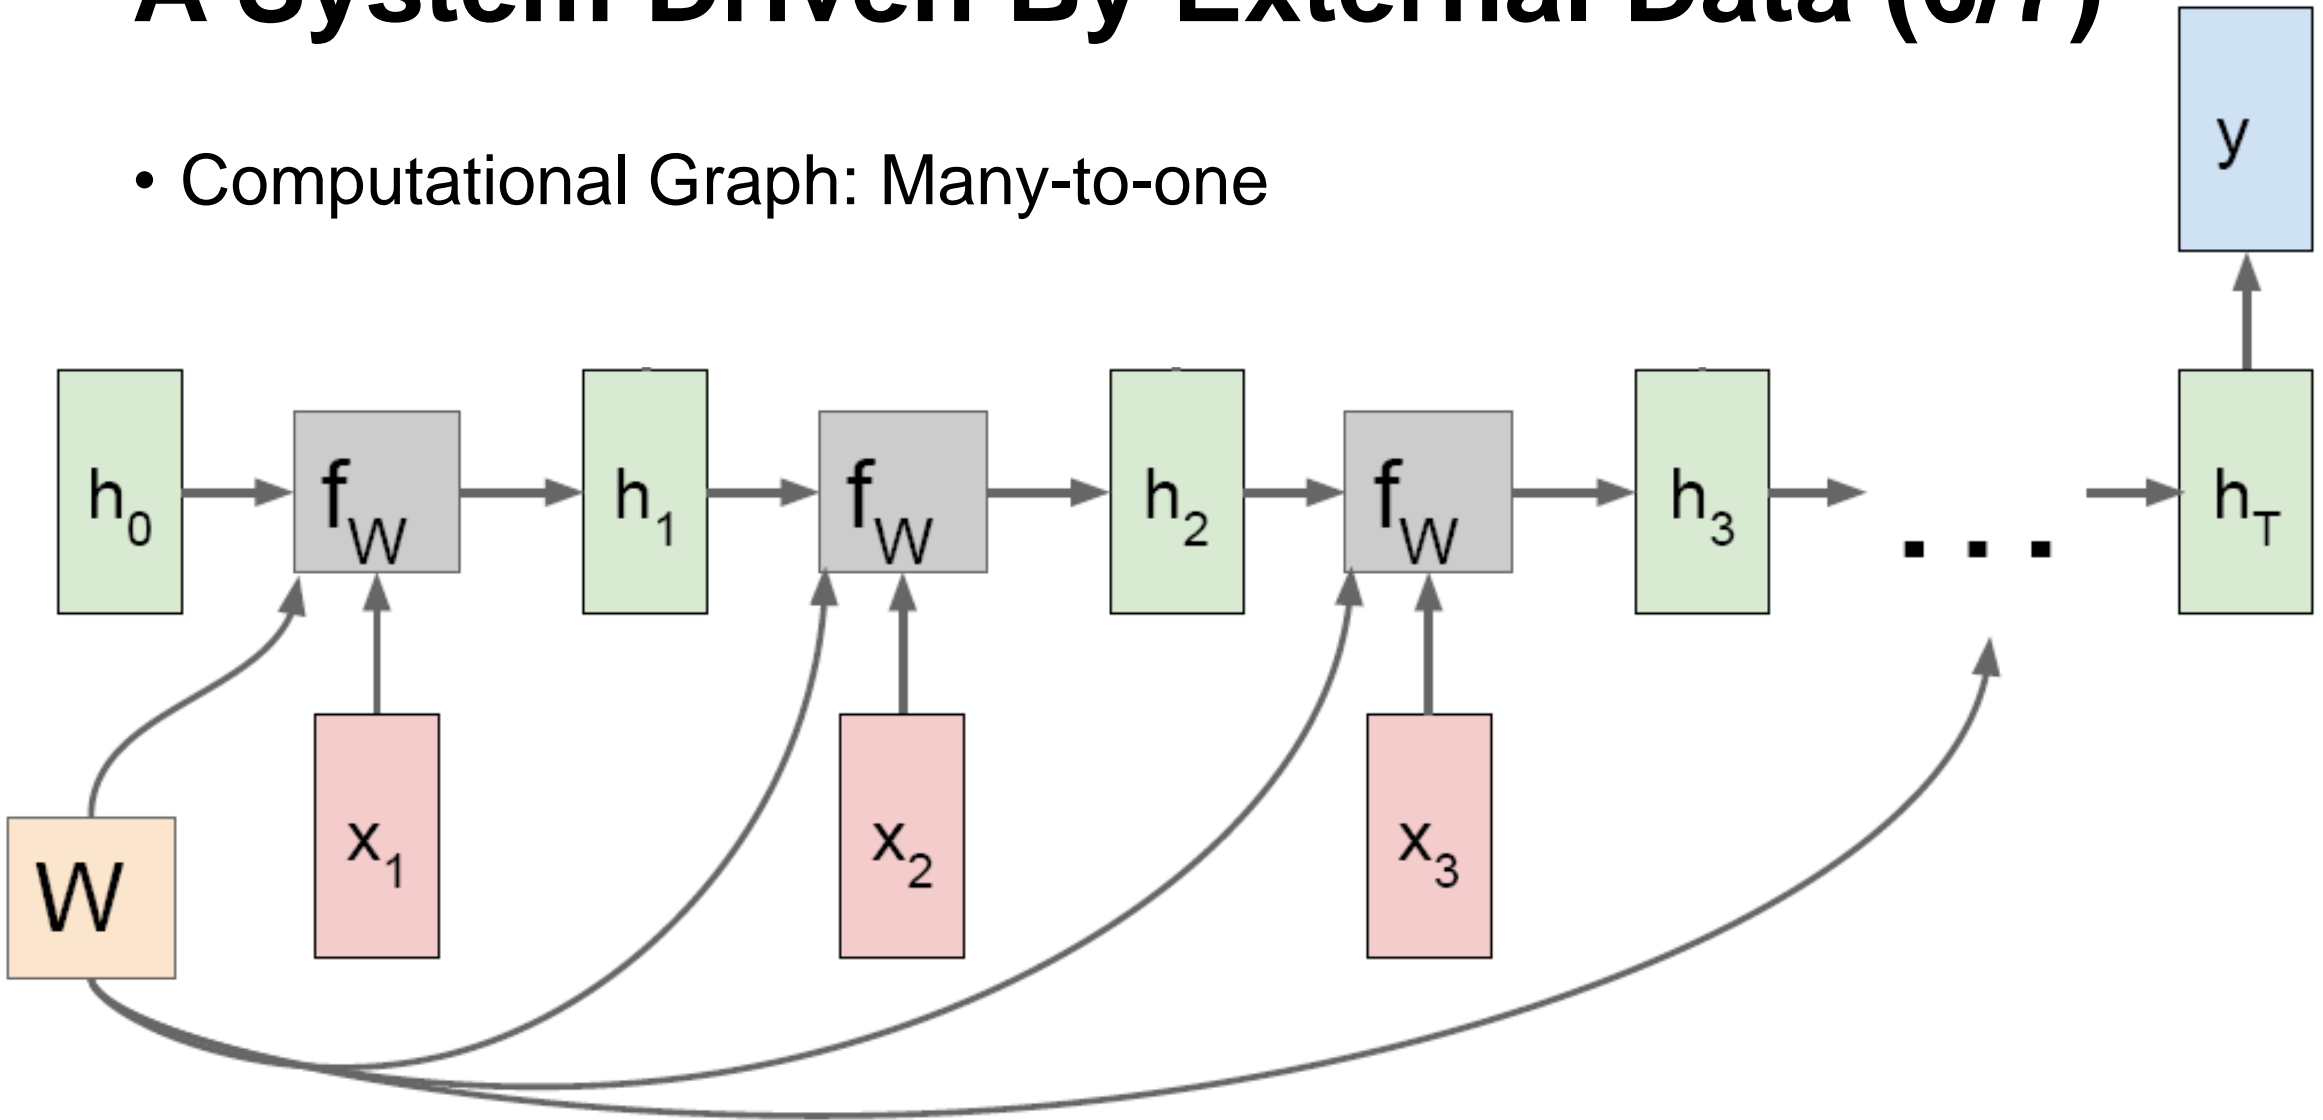
A System Driven By External Data (5/7)

- Computational Graph:
Many-to-many



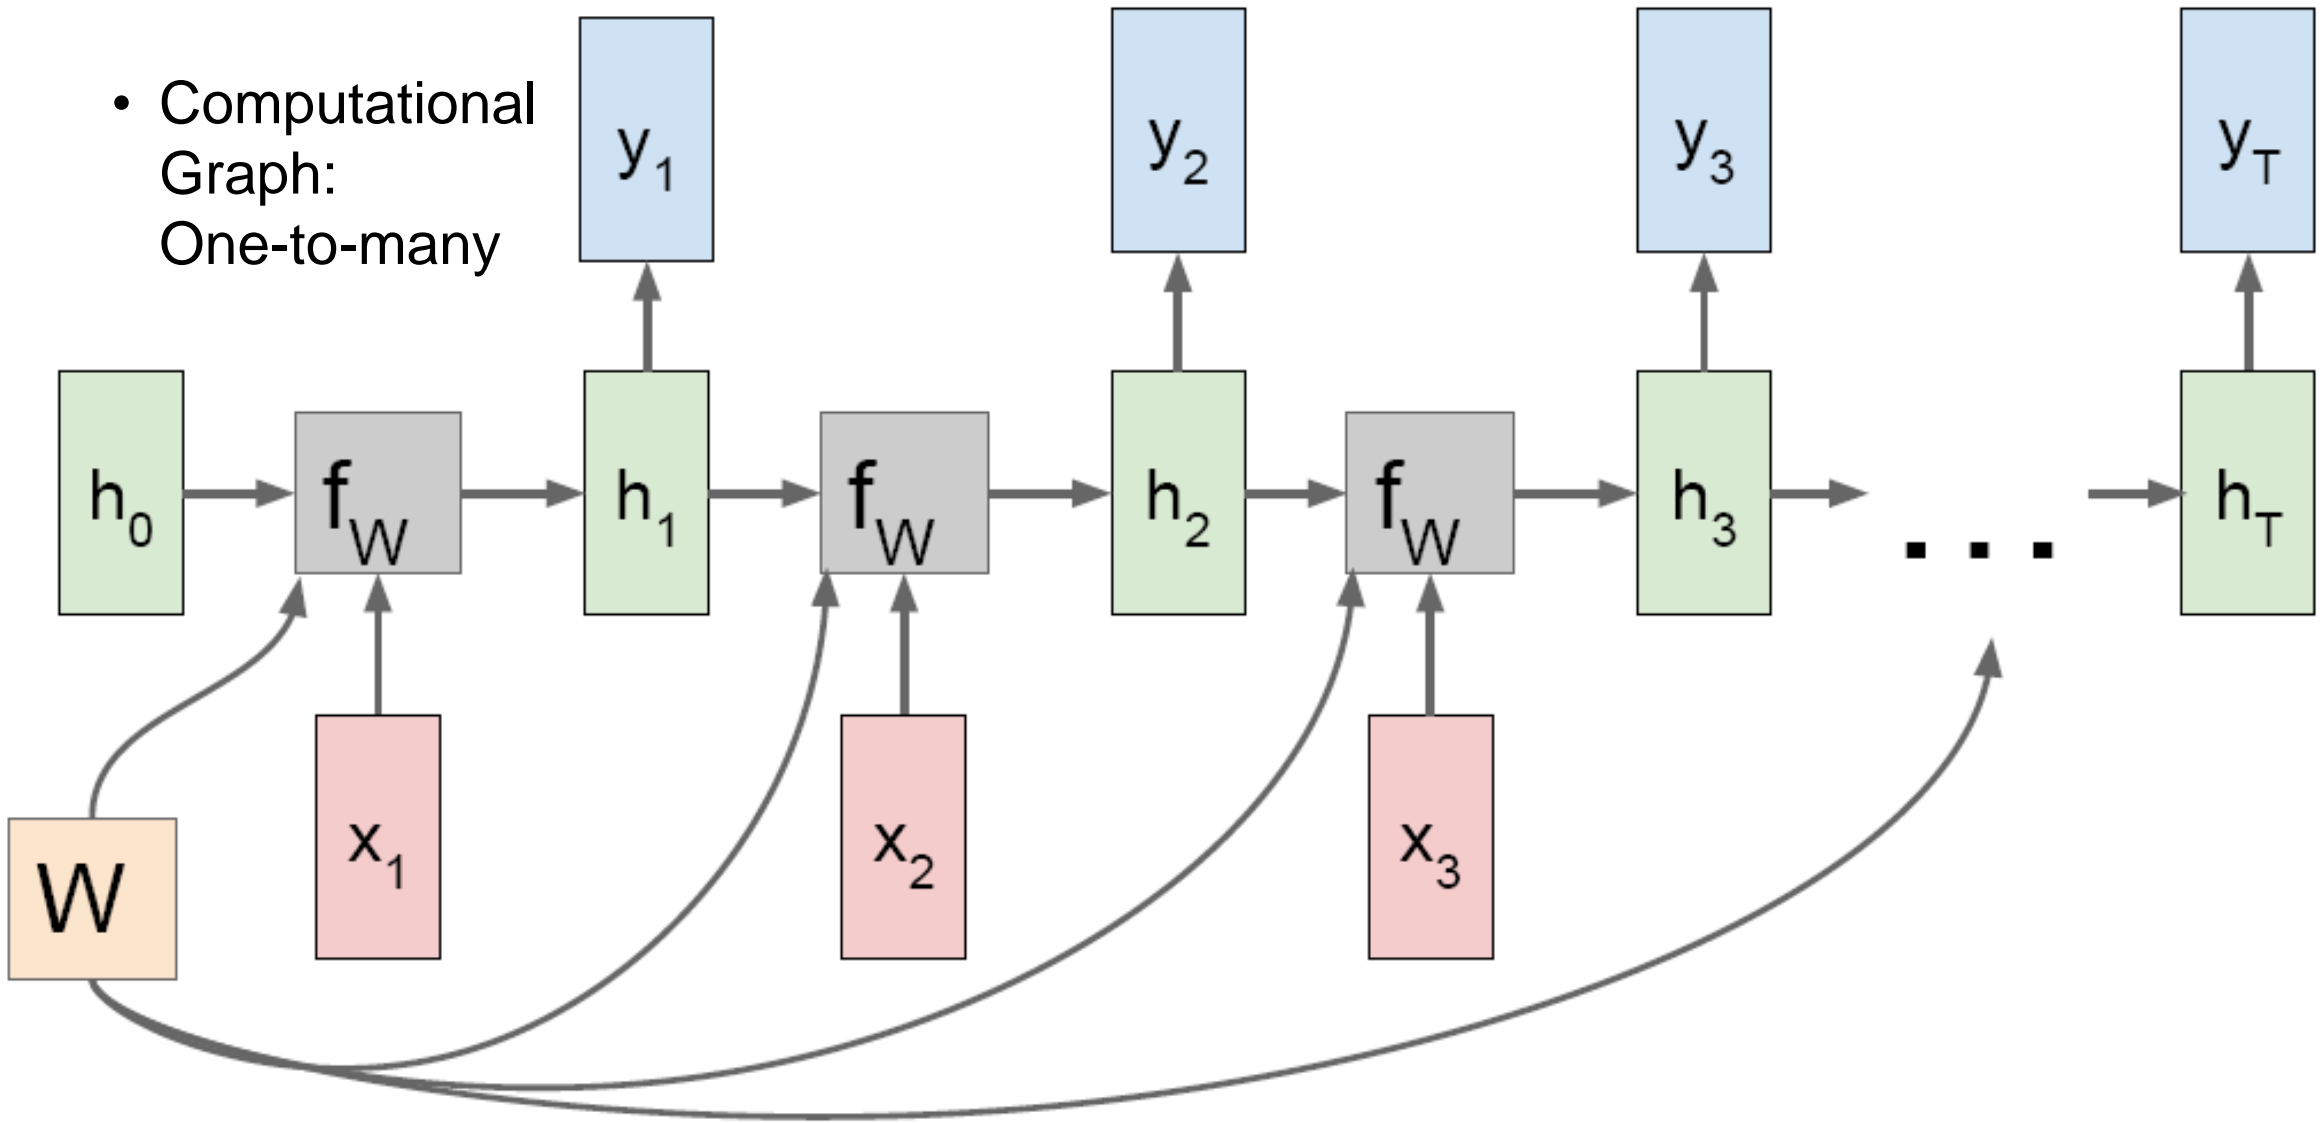
A System Driven By External Data (6/7)

- Computational Graph: Many-to-one



A System Driven By External Data (7/7)

- Computational Graph:
One-to-many



Recurrent Neural Networks (1/8)

- Use **the same** computational function and parameters across different time steps of the sequence
- Each time step: takes the input entry and **the previous hidden state** to compute the output entry
- Loss: typically computed at every time step

Recurrent Neural Networks (2/8)

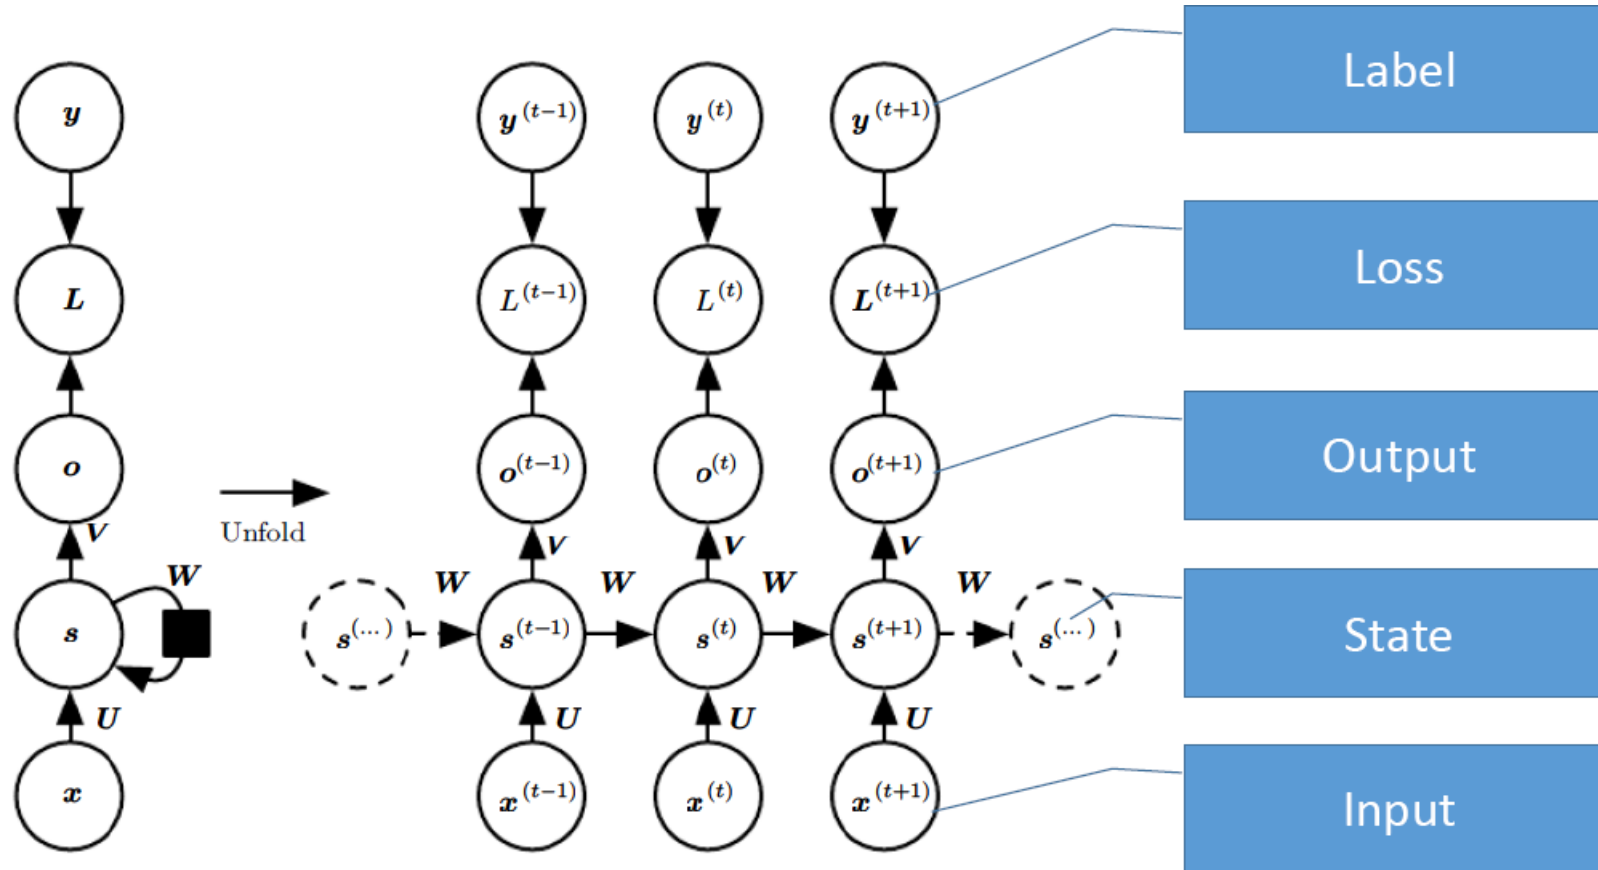


Figure from Deep Learning, Goodfellow, Bengio and Courville

Recurrent Neural Networks (3/8)

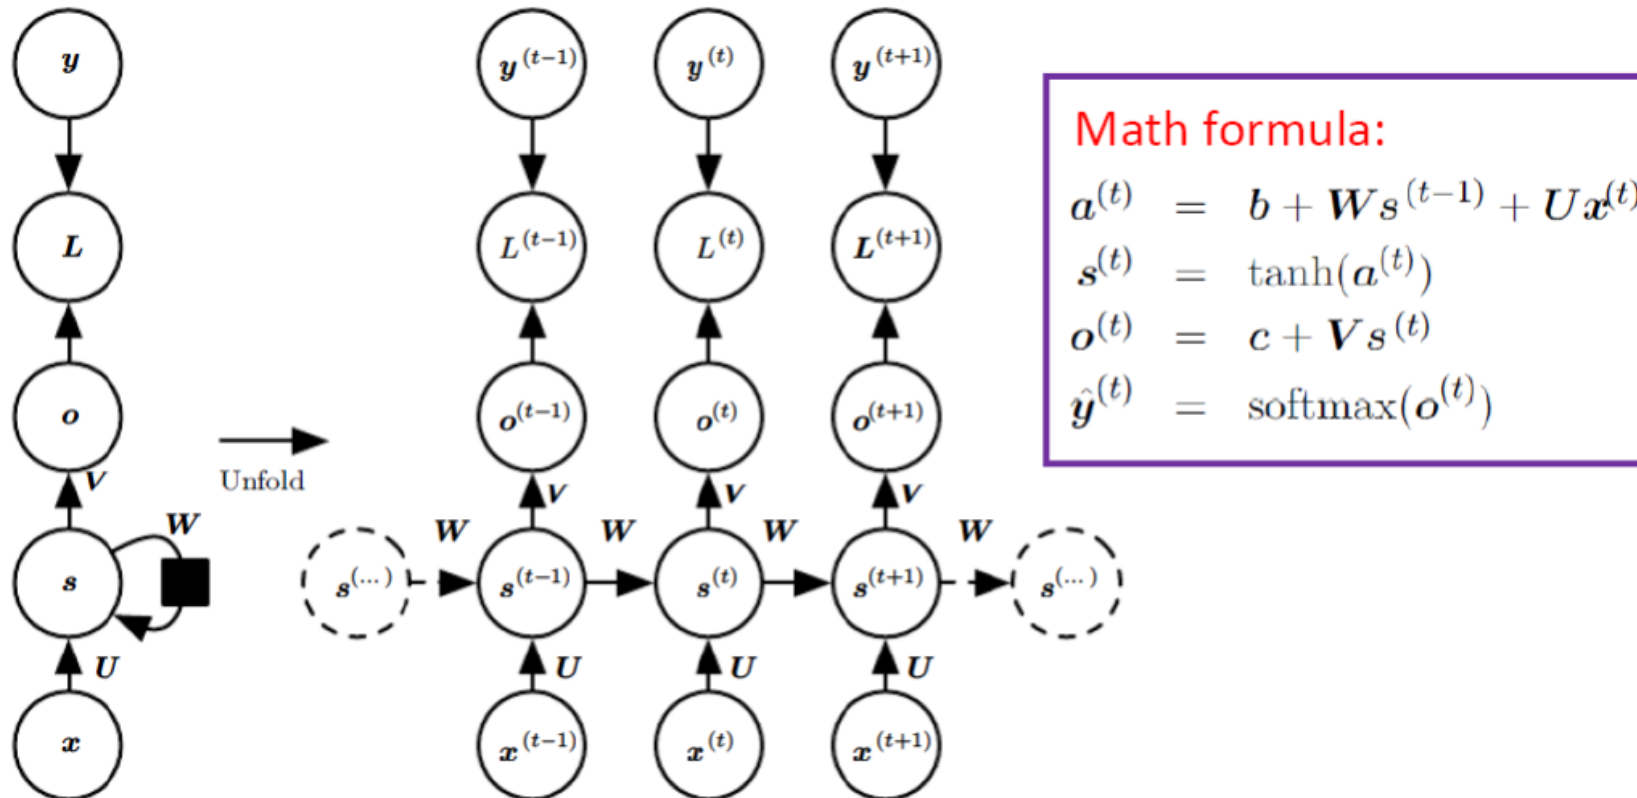



Figure from Deep Learning, Goodfellow, Bengio and Courville

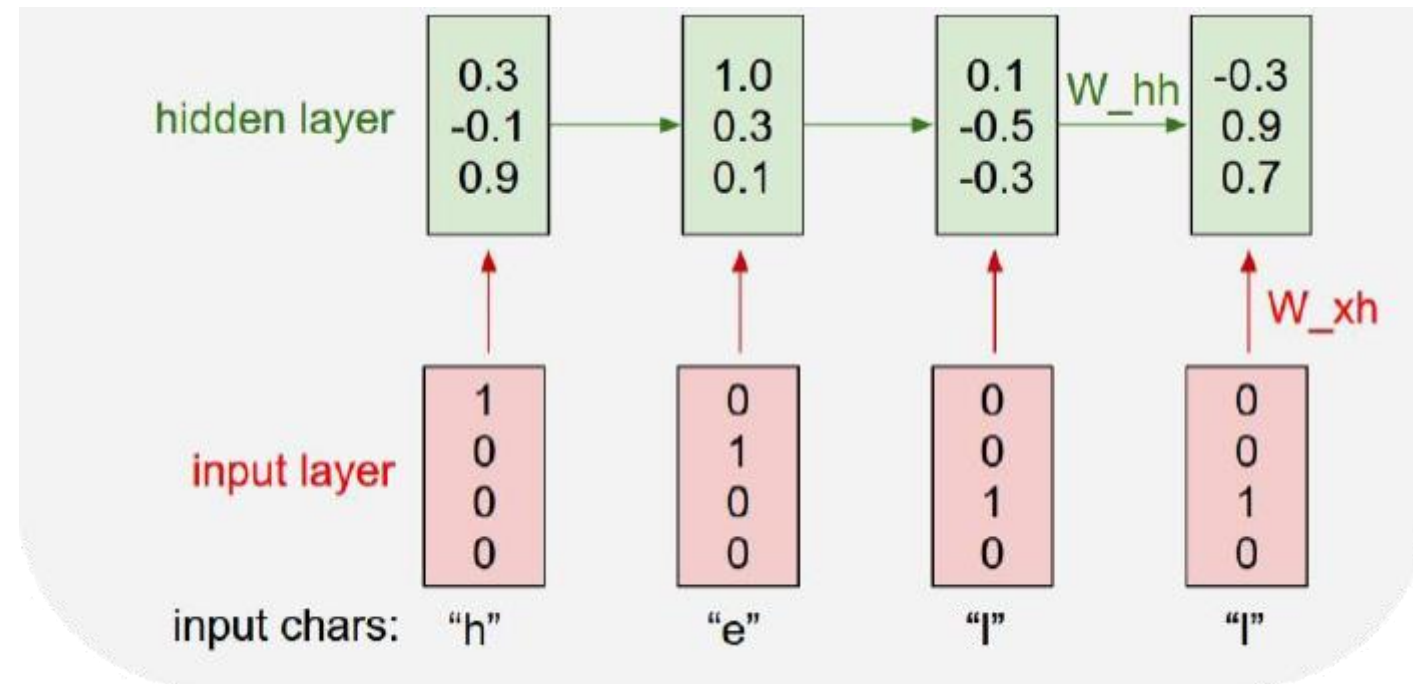
Recurrent Neural Networks (4/8)

- **Example:**
 - Character-level language model
 - Vocabulary: [h,e,l,o]
 - Training sample: “hello”



Recurrent Neural Networks (5/8)

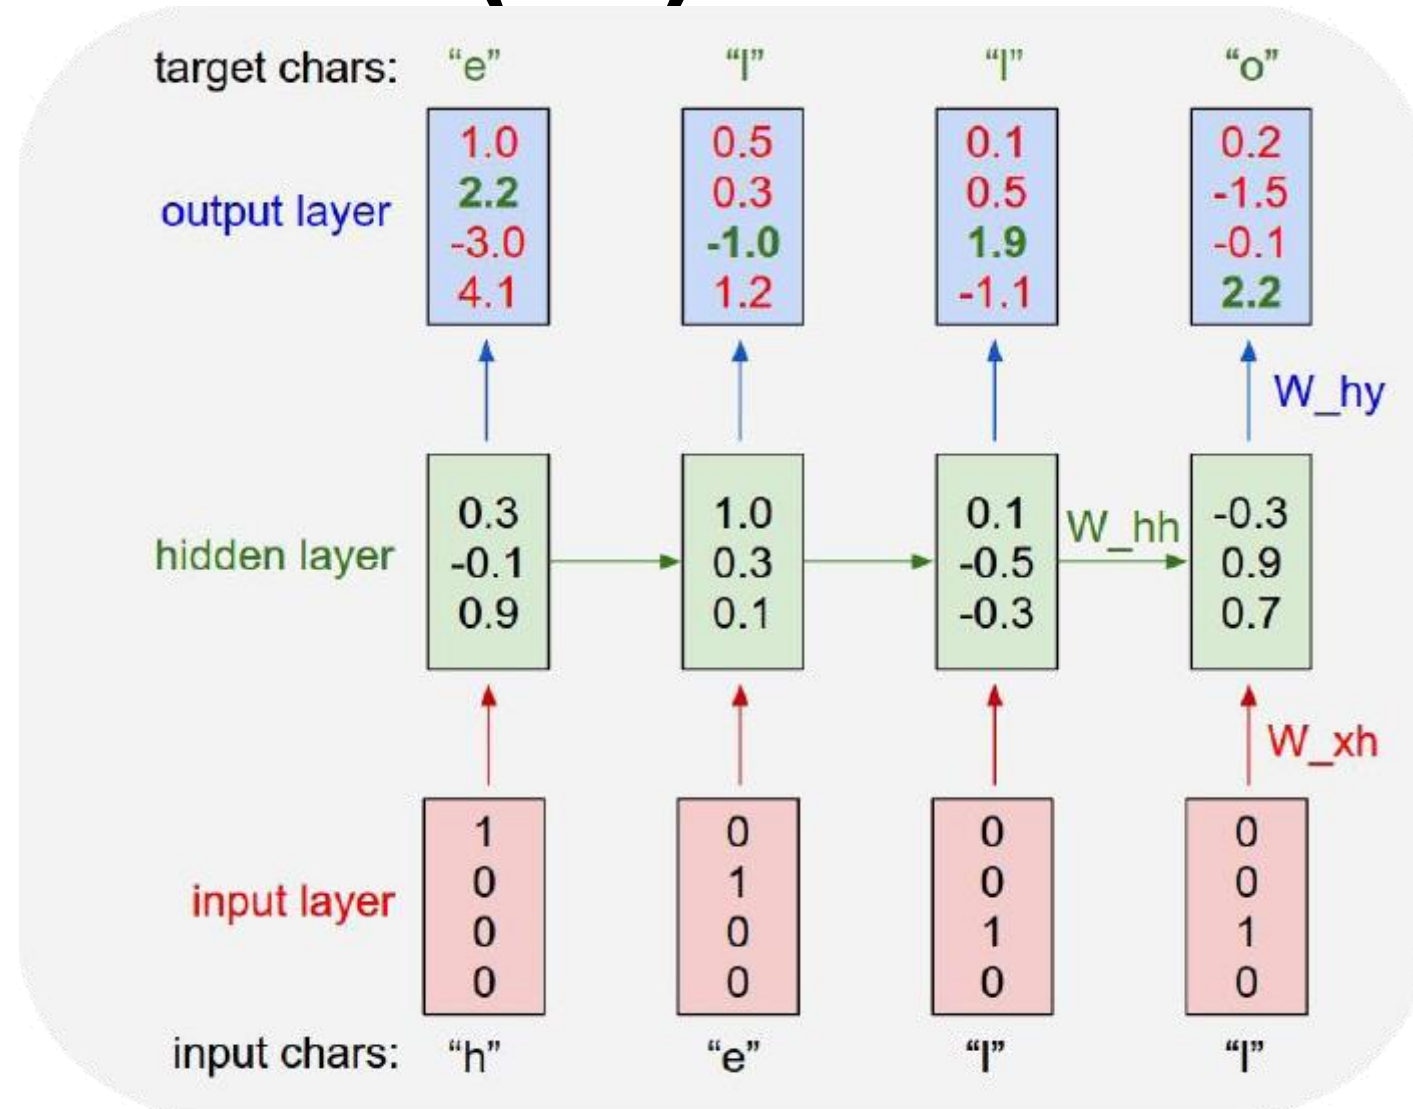
- **Example:**
 - Character-level language model
 - Vocabulary: [h,e,l,o]
 - Training sample: “hello”



Recurrent Neural Networks (6/8)

- **Example:**

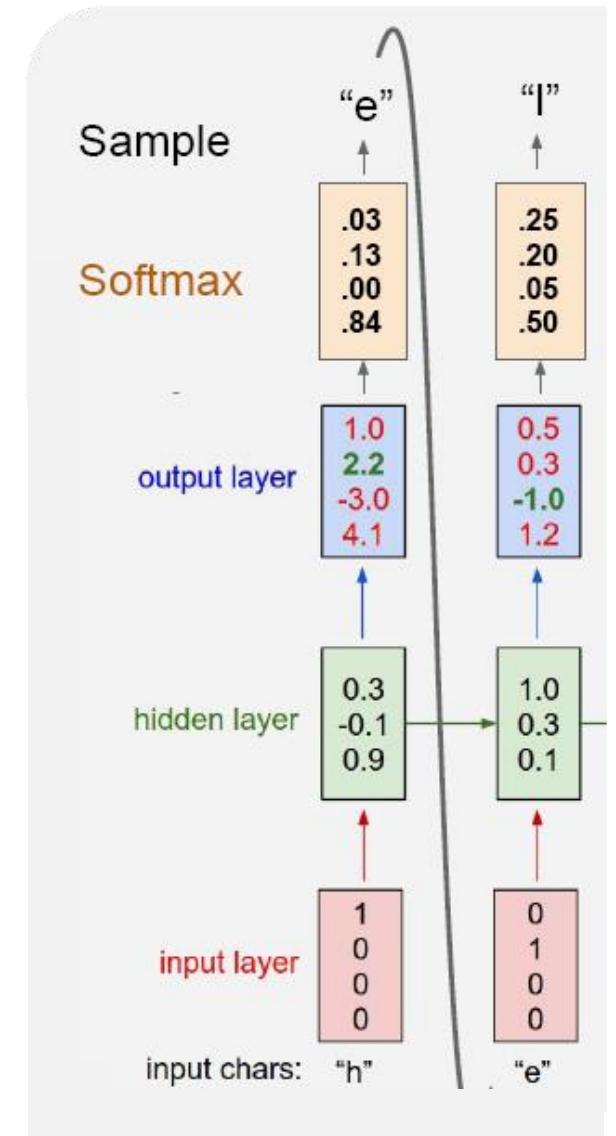
- Character-level language model
- Vocabulary: [h,e,l,o]
- Training sample: "hello"



Recurrent Neural Networks (7/8)

- **Example:**

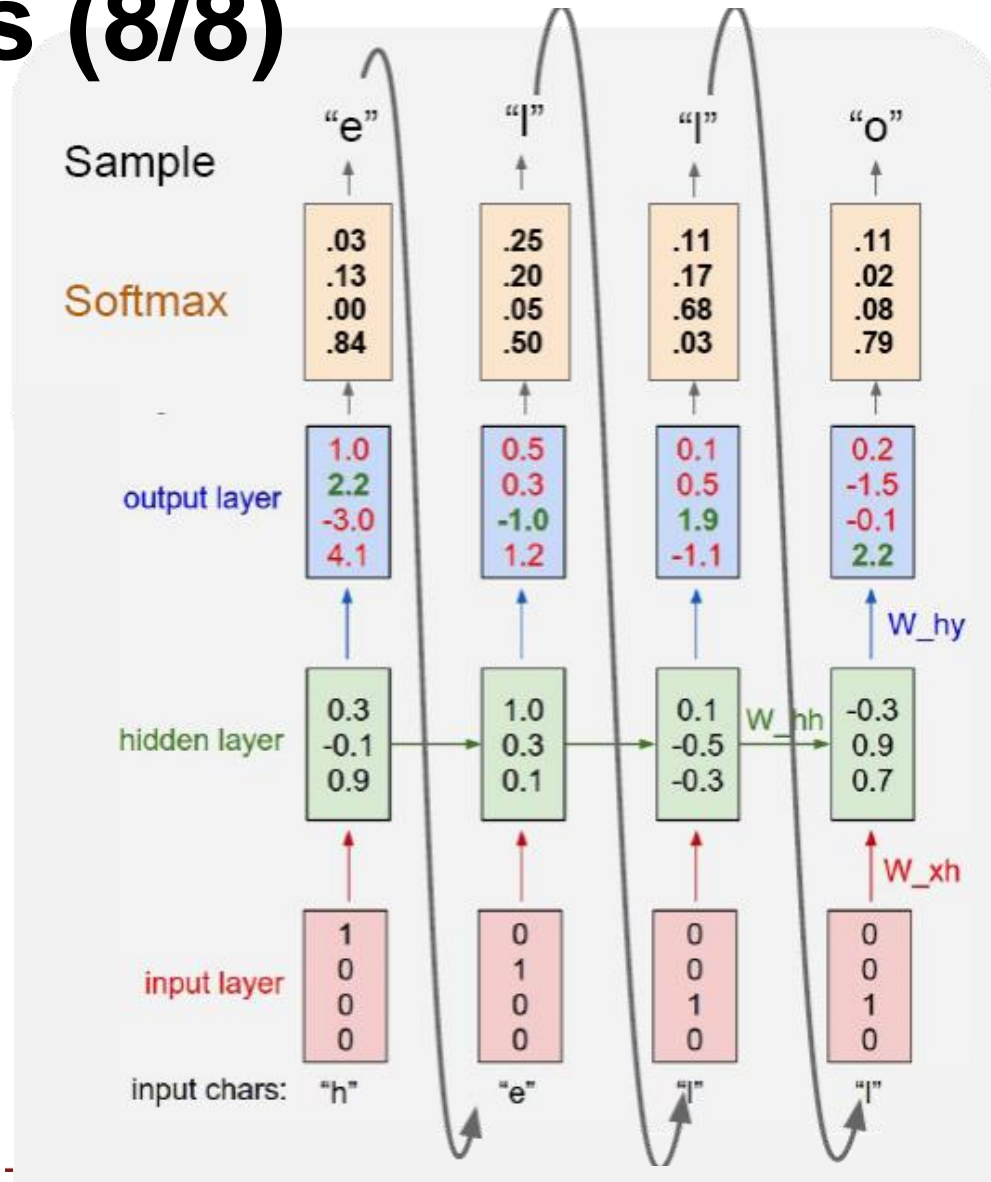
- Character-level language model
- Vocabulary: [h,e,l,o]
- Testing: character one at a time, feed back to model



Recurrent Neural Networks (8/8)

- **Example:**

- Character-level language model
- Vocabulary: [h,e,l,o]
- Testing: character one at a time, feed back to model



Advantages

- **Hidden state**: a lossy summary of the past
- Shared functions and parameters: greatly reduce the capacity and good for **generalization** in learning
- Explicitly use the **prior knowledge** that the sequential data can be processed by in the same way at different time step (e.g., NLP)
- Yet still powerful (actually **universal**): any function computable by a Turing machine can be computed by such a recurrent network of a finite size (see, e.g., Siegelmann and Sontag [1995])

Example Implementation

- Given past, predict future; let's implement an MLP model

LSTM: Long Short Term Memory

- The basic structure of LSTM and some symbols to aid understanding

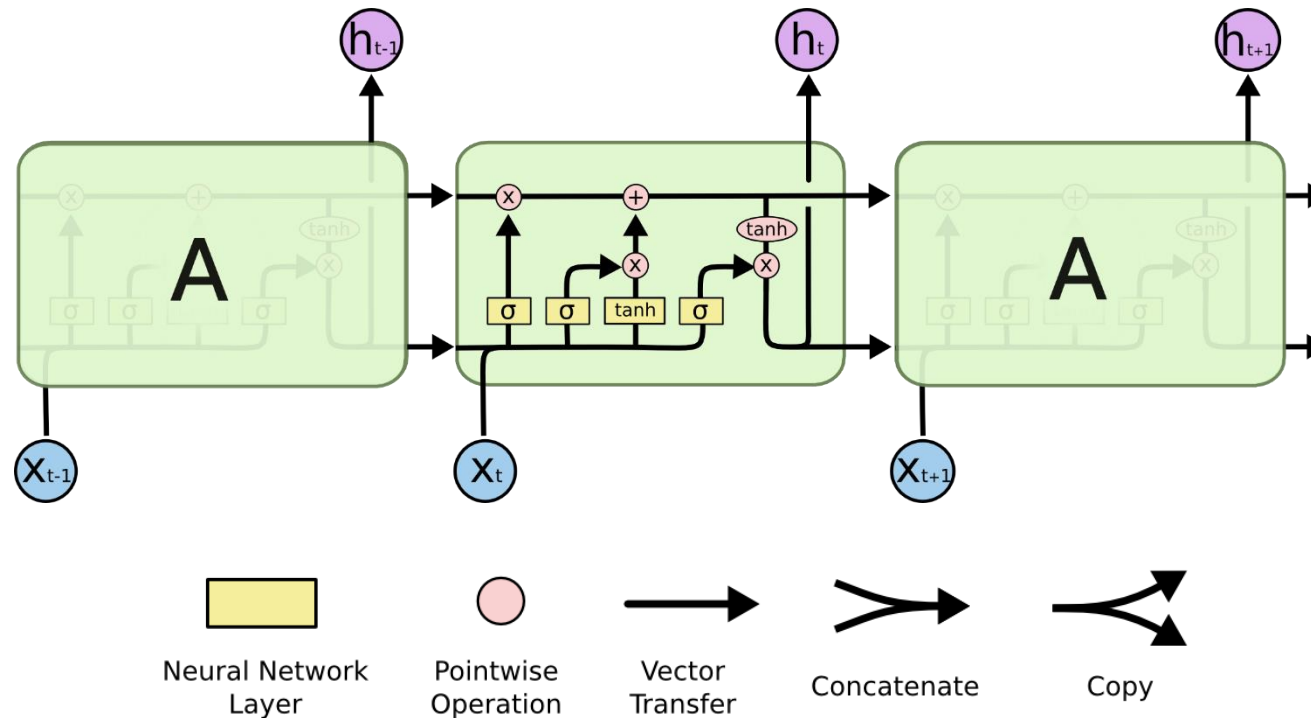



Image source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs>

LSTM Core Ideas

- Two key ideas of LSTM:
 - A backbone to carry state forward and gradients backward.



- Gating (pointwise multiplication) to modulate information flow. Sigmoid makes $0 < \text{gate} < 1$.





Image source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs>

LSTM Gating: Forget

- The f gate is ‘forgetting.’ Use previous state, C , previous output, h , and current input, x , to determine how much to suppress previous state.
- E.g., C might encode the fact that we have a subject and need a verb. Forget that when verb found.




$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

Image source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs>

LSTM Gating: Input Gate

- Input gate i determines which values of C to update
- Separate tanh layer produces new state to add to C




$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Image source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs>

LSTM Gating: Update to C

- Forget gate does pointwise modulation of C.
- Input gate modulates the tanh layer – this is added to C.




$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Image source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs>

LSTM Gating: Output

- o is the **output gate**: modulates what part of the state C gets passed (via \tanh) to current output h



- E.g., could encode whether a noun is singular or plural to prepare for a verb
- But the real features are learned, not engineered.


$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Image source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs>

The Popular LSTM Cell

- Another demonstration




**Dashed line
indicates time-lag**

Example Implementation

- Given past, predict future; let's implement a simple LSTM model
- Discuss pros/cons, compare with MLP model

GRU: Gated Recurrent Unit

- Combine C and h into a single state/output
- Combine forget and input gates into update gate, z



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Image source: colah.github.io/posts/2015-08-Understanding-LSTMs

Gated Recurrent Units (GRUs) (1/3)

- Main idea:
 - Keep around memory to capture **long dependencies**
 - Allow error messages to flow at **different strengths** depending on the inputs
- Standard RNN computes hidden layer at next time step directly: $h_t = \sigma(W^{(hh)}h_{t-1} + W^{(hx)}x_t)$
- Compute an update gate based on current input word vector and hidden state $z_t = \sigma(U^{(z)}h_{t-1} + W^{(z)}x_t)$
 - Controls how much of past state should matter now
 - If z close to 1, then we can copy information in that unit through many steps!




Image source: www.wildml.com/2015/10/recurrent-neural-network-tutorial-part-4-implementing-a-grulstm-rnn-with-python-and-theano

Gated Recurrent Units (GRUs) (2/3)

- Standard RNN computes hidden layer at next time step directly: $h_t = \sigma(W^{(hh)}h_{t-1} + W^{(hx)}x_t)$
- Compute an update gate based on current input word vector and hidden state
$$z_t = \sigma(U^{(z)}h_{t-1} + W^{(z)}x_t)$$
- Compute a reset gate similarly but with different weights $r_t = \sigma(U^{(r)}h_{t-1} + W^{(r)}x_t)$

- Units with **short-term** dependencies often have **reset** gates very active
- Units with **long-term** dependencies have active **update** gates z

- Standard RNN computes hidden layer at next time step directly: $h_t = \sigma(W^{(hh)}h_{t-1} + W^{(hx)}x_t)$
- Compute an update gate based on current input word vector and hidden state
$$z_t = \sigma(U^{(z)}h_{t-1} + W^{(z)}x_t)$$
- Compute a reset gate similarly but with different weights $r_t = \sigma(U^{(r)}h_{t-1} + W^{(r)}x_t)$
 - Units with **short-term** dependencies often have **reset** gates very active
 - Units with **long-term** dependencies have active **update** gates z

Image source: www.wildml.com/2015/10/recurrent-neural-network-tutorial-part-4-implementing-a-grulstm-rnn-with-python-and-theano

If reset close to 0, ignore previous hidden state (allows model to drop information that is irrelevant in the future)

Gated Recurrent Units (GRUs) (3/3)

- Standard RNN computes hidden layer at next time step directly $h_t = \sigma(W^{(hh)}h_{t-1} + W^{(hx)}x_t)$
- Compute an update gate based on current input word vector and hidden state
$$z_t = \sigma(U^{(z)}h_{t-1} + W^{(z)}x_t)$$
- Compute a reset gate similarly but with different weights $r_t = \sigma(U^{(r)}h_{t-1} + W^{(r)}x_t)$
- New memory $\tilde{h}_t = \tanh(r_t \circ U h_{t-1} + W x_t)$
- Final memory $h_t = z_t \circ h_{t-1} + (1 - z_t) \circ \tilde{h}_t$

- Standard RNN computes hidden layer at next time step directly $h_t = \sigma(W^{(hh)}h_{t-1} + W^{(hx)}x_t)$
- Compute an update gate based on current input word vector and hidden state
$$z_t = \sigma(U^{(z)}h_{t-1} + W^{(z)}x_t)$$
- Compute a reset gate similarly but with different weights $r_t = \sigma(U^{(r)}h_{t-1} + W^{(r)}x_t)$
- New memory $\tilde{h}_t = \tanh(r_t \circ U h_{t-1} + W x_t)$
- Final memory $h_t = z_t \circ h_{t-1} + (1 - z_t) \circ \tilde{h}_t$

Image source: www.wildml.com/2015/10/recurrent-neural-network-tutorial-part-4-implementing-a-grulstm-rnn-with-python-and-theano

- LSTMs are a more complex form, but basically same intuition
- GRUs are often more preferred than LSTMs

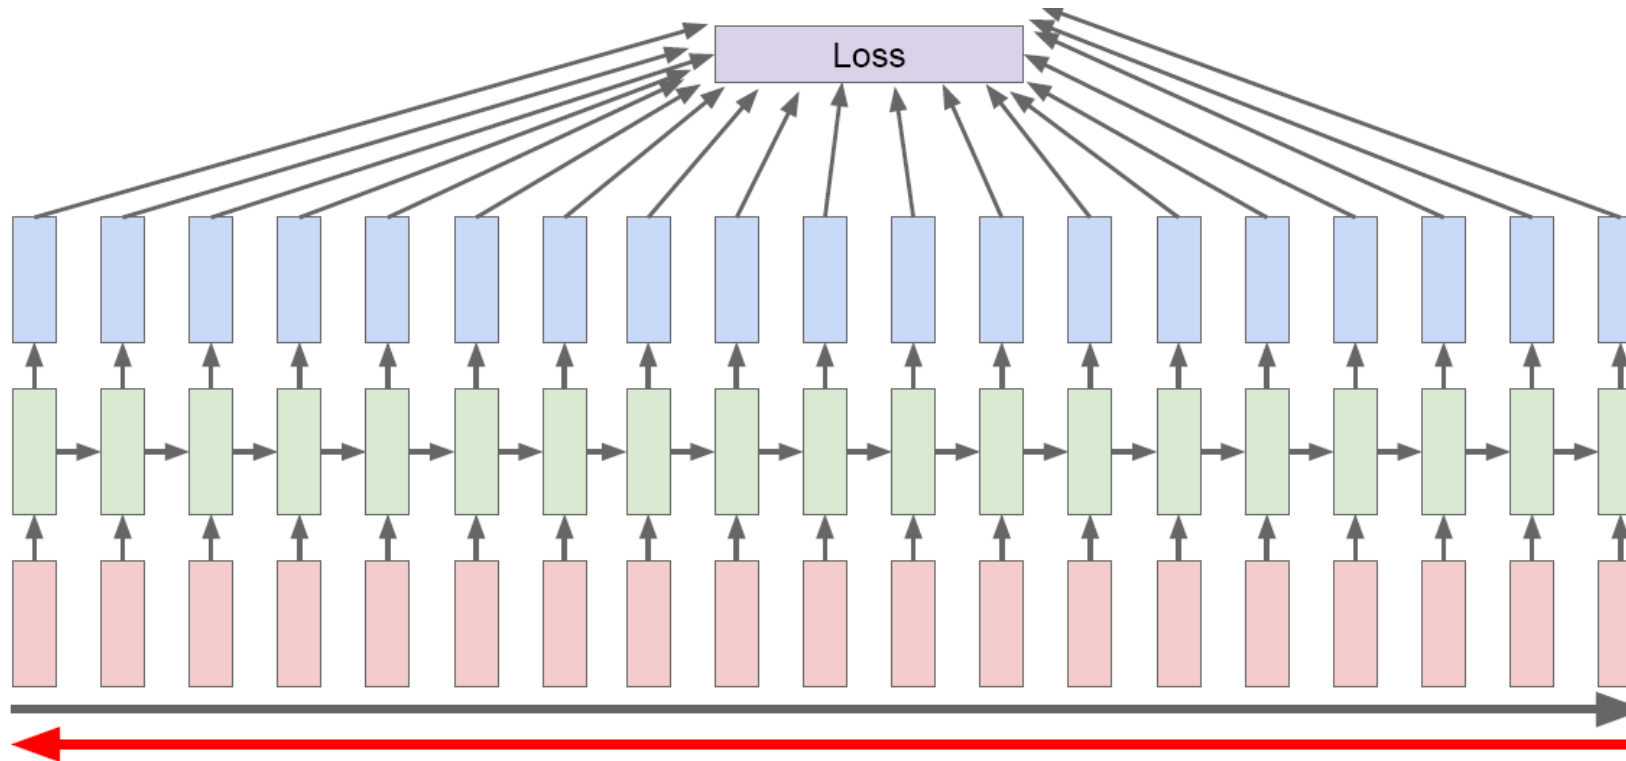
combines current & previous time steps

Training RNN (1/7)

- Principle: unfold the computational graph, and use backpropagation
- Called **back-propagation through time (BPTT) algorithm**
- Can then apply any general-purpose gradient-based techniques
- Conceptually: first compute the gradients of **the internal nodes**, then compute the gradients of **the parameters**

Training RNN (2/7)

- Backpropagation through time



Training RNN (3/7)

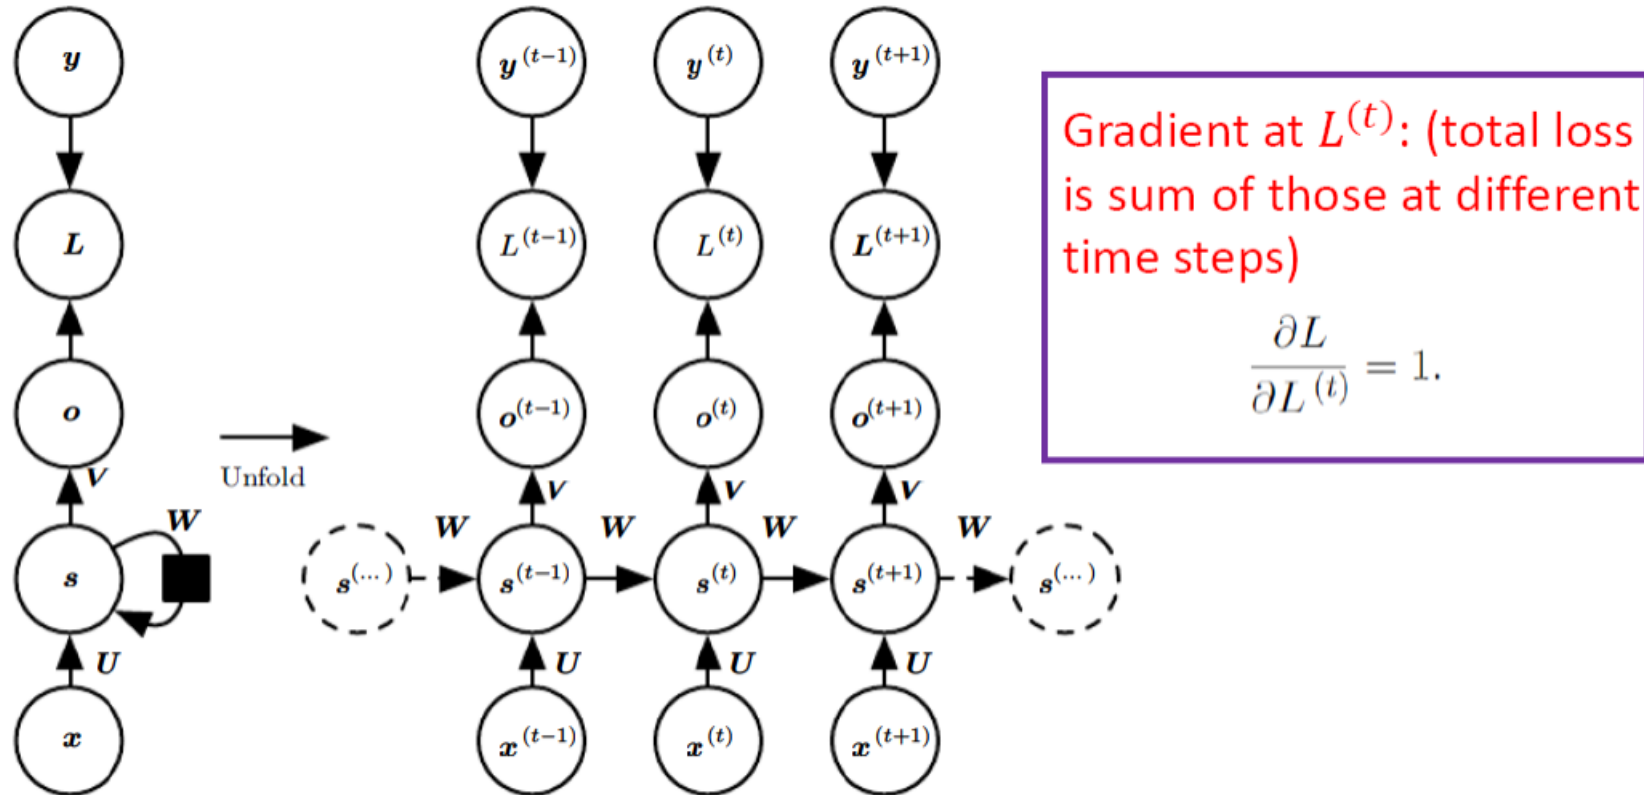


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

Training RNN (4/7)

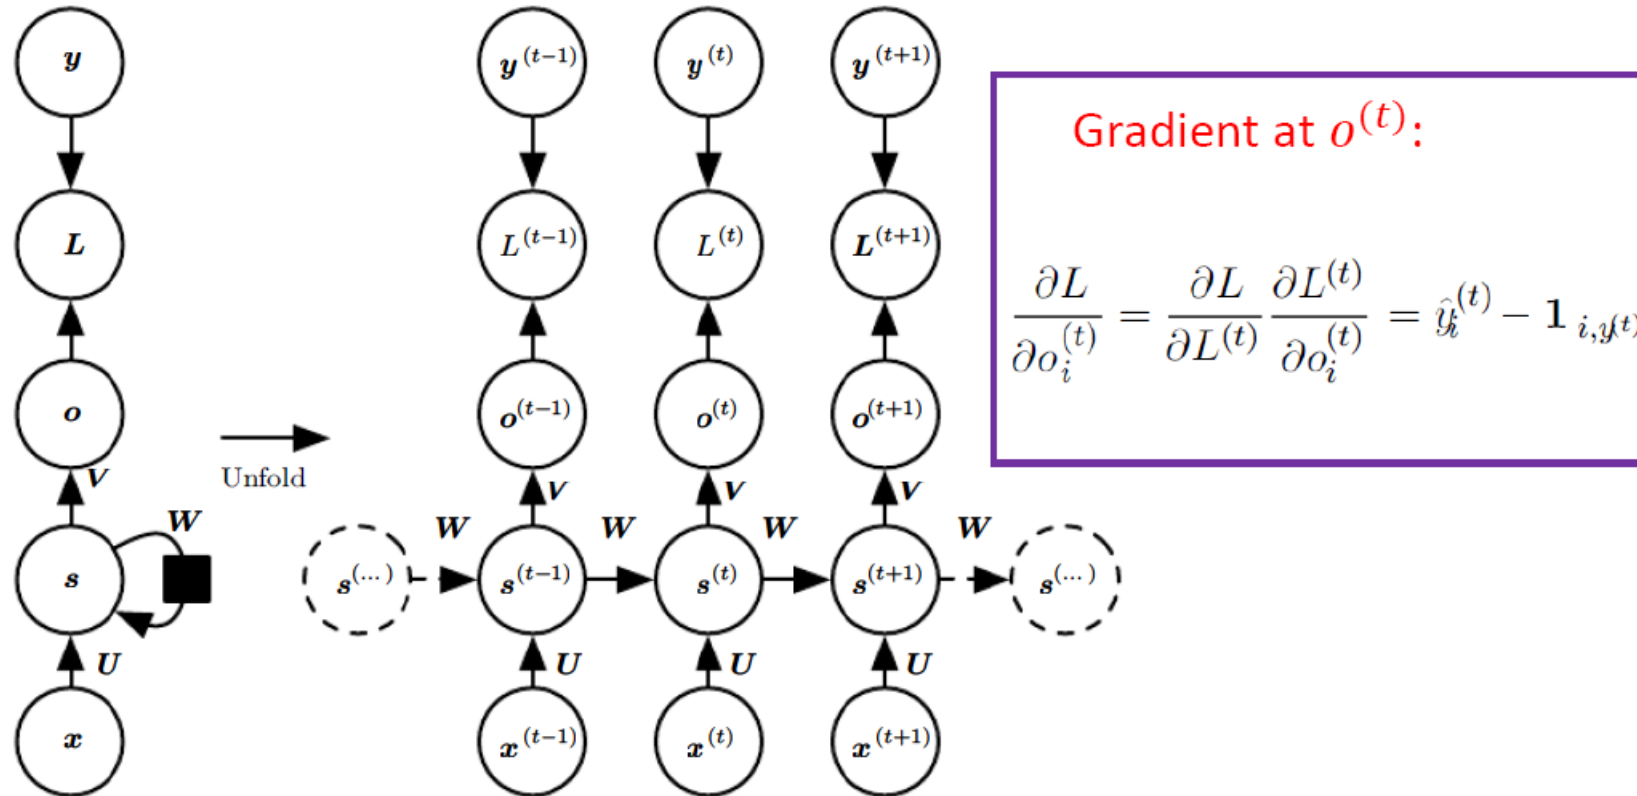


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

Training RNN (5/7)

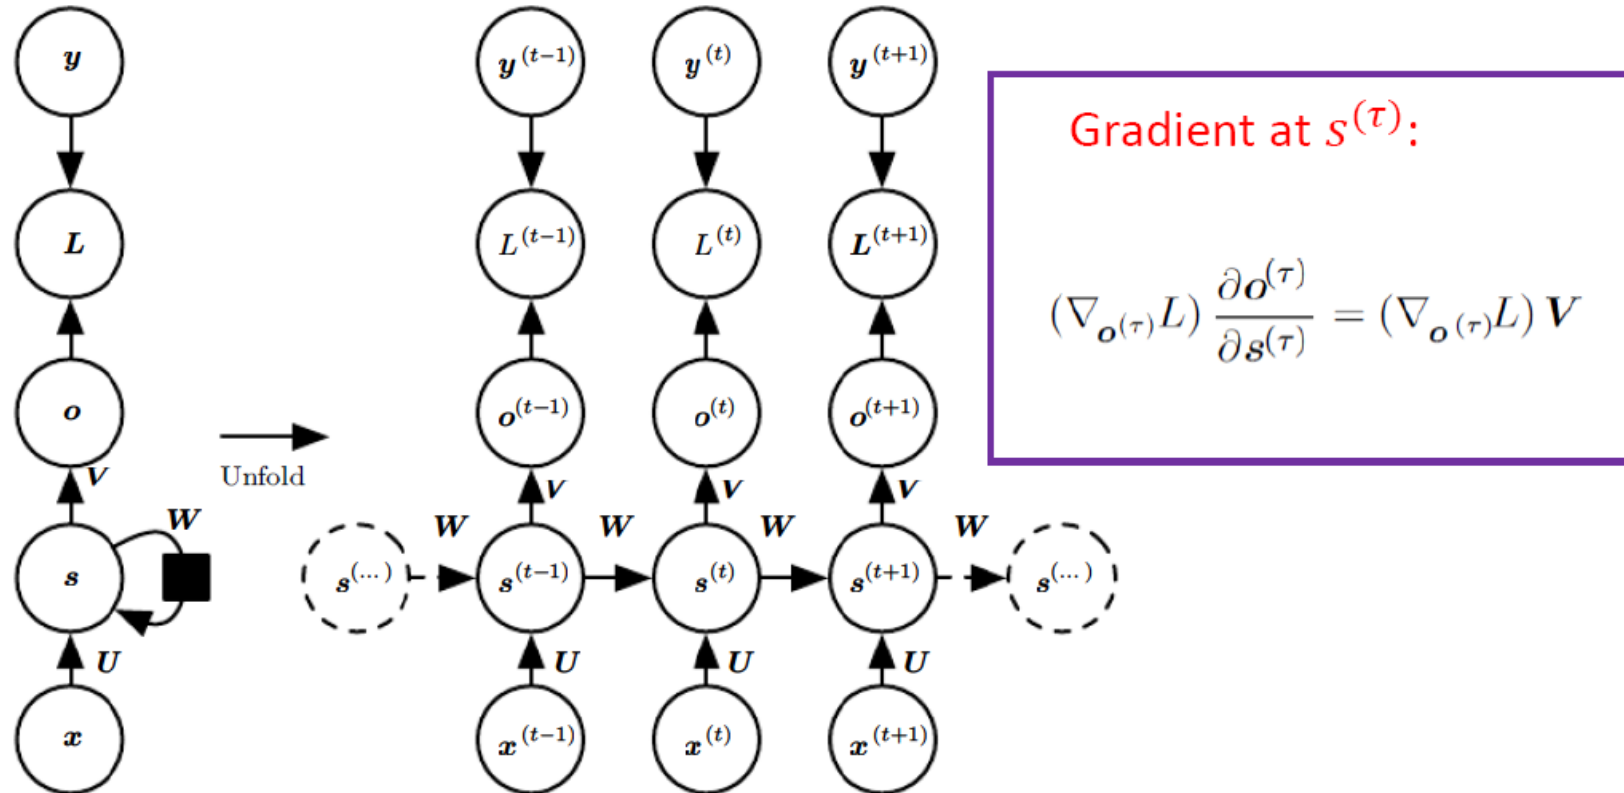


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

Training RNN (6/7)

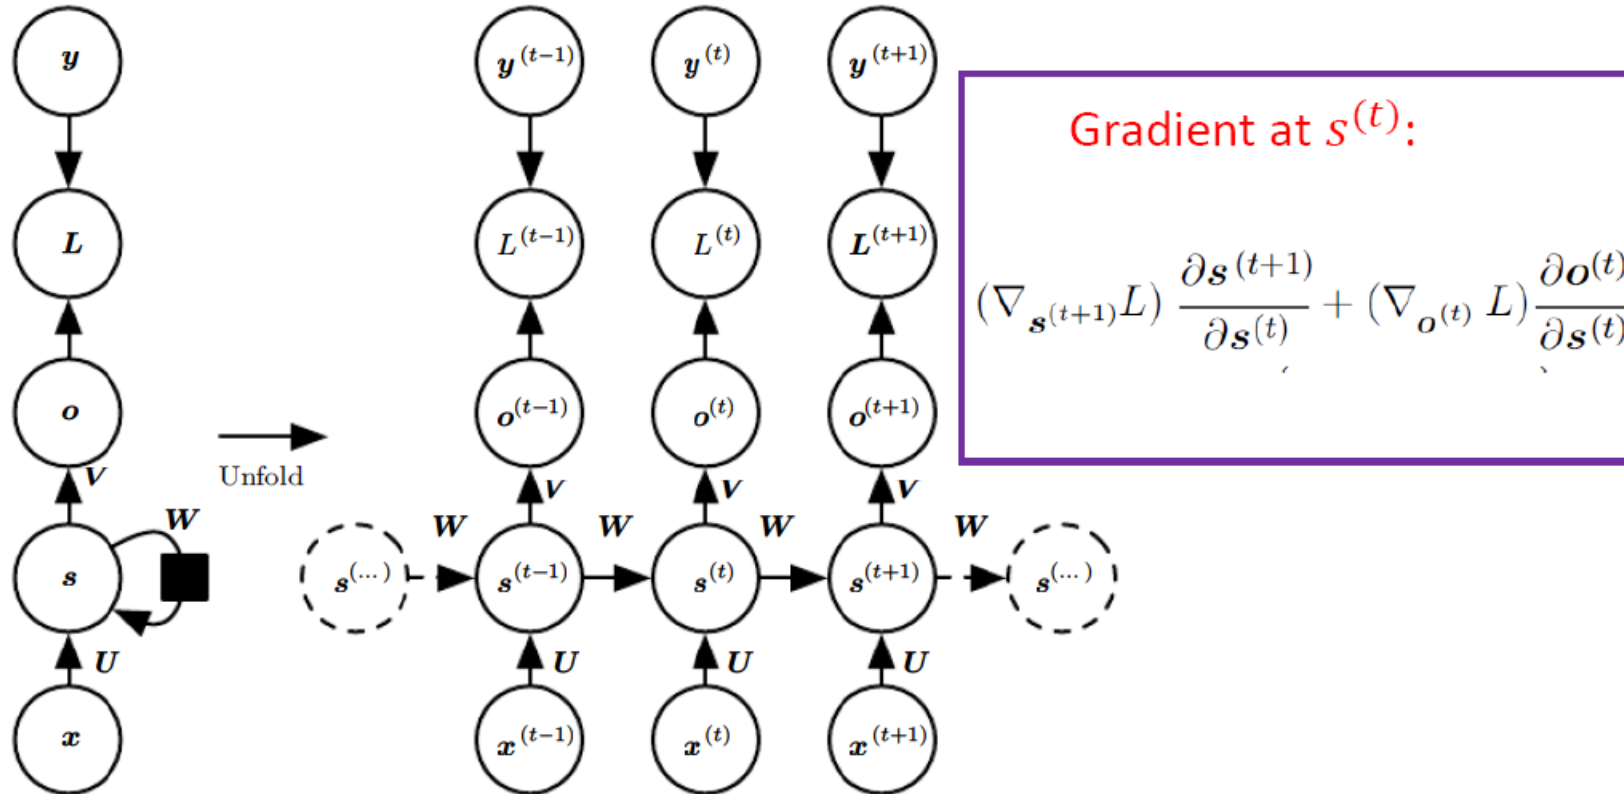


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

Training RNN (7/7)

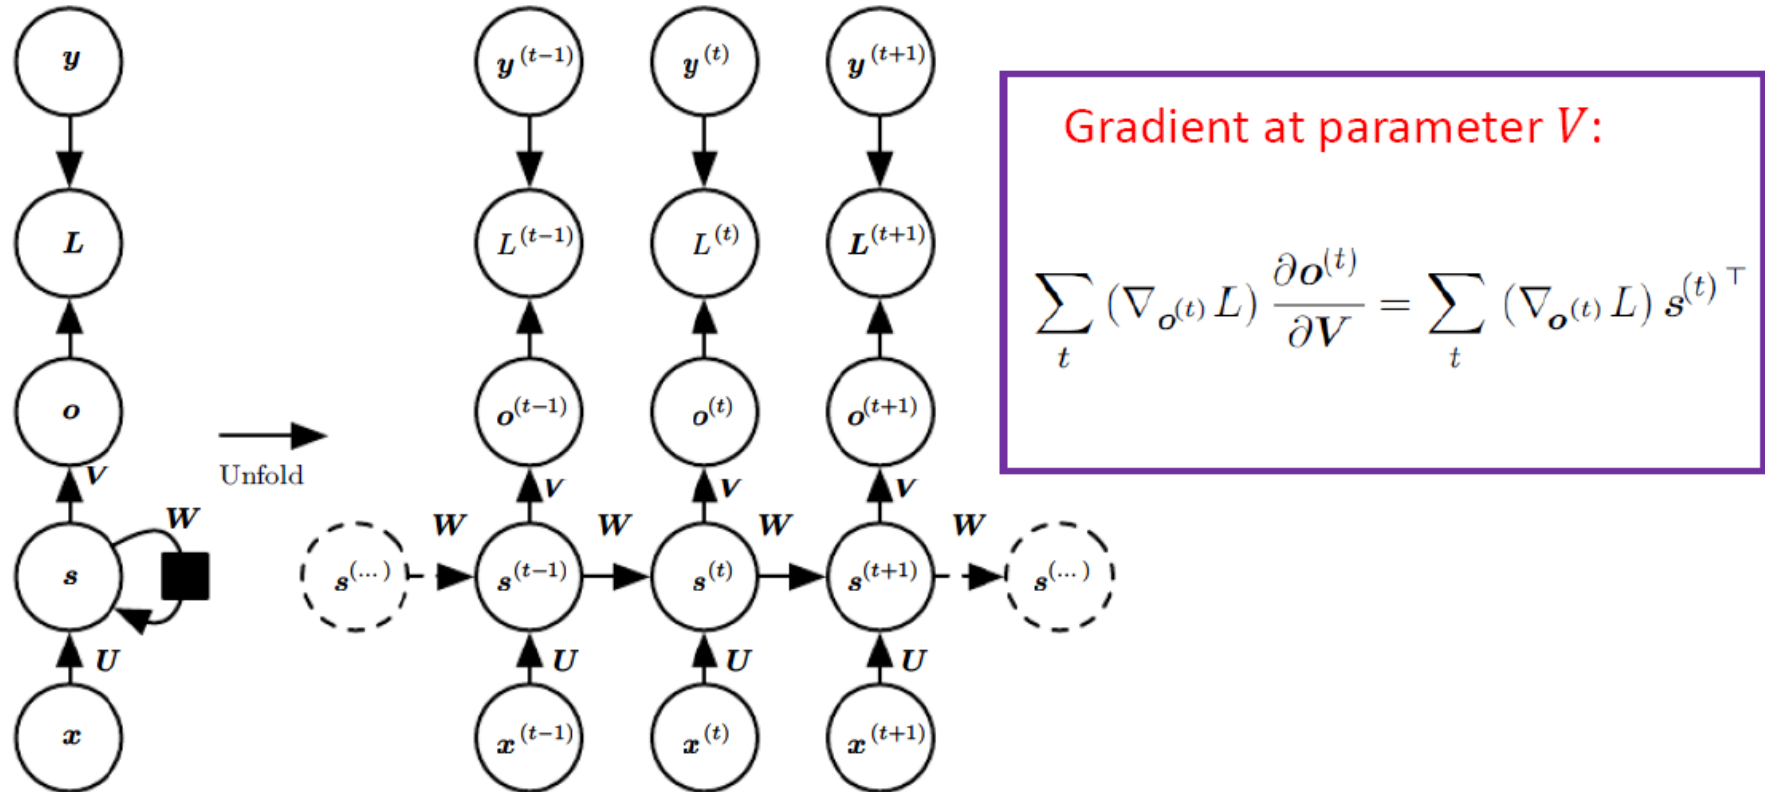


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

Recurrent Neural Networks

- Use **the same** computational function and parameters across different time steps of the sequence
- Each time step: takes the input entry and **the previous hidden state** to compute the output entry
- Loss: typically computed at every time step
- Many variants
 - Information about the past can be in many other forms
 - Only output at the end of the sequence

RNN Variations (1/2)

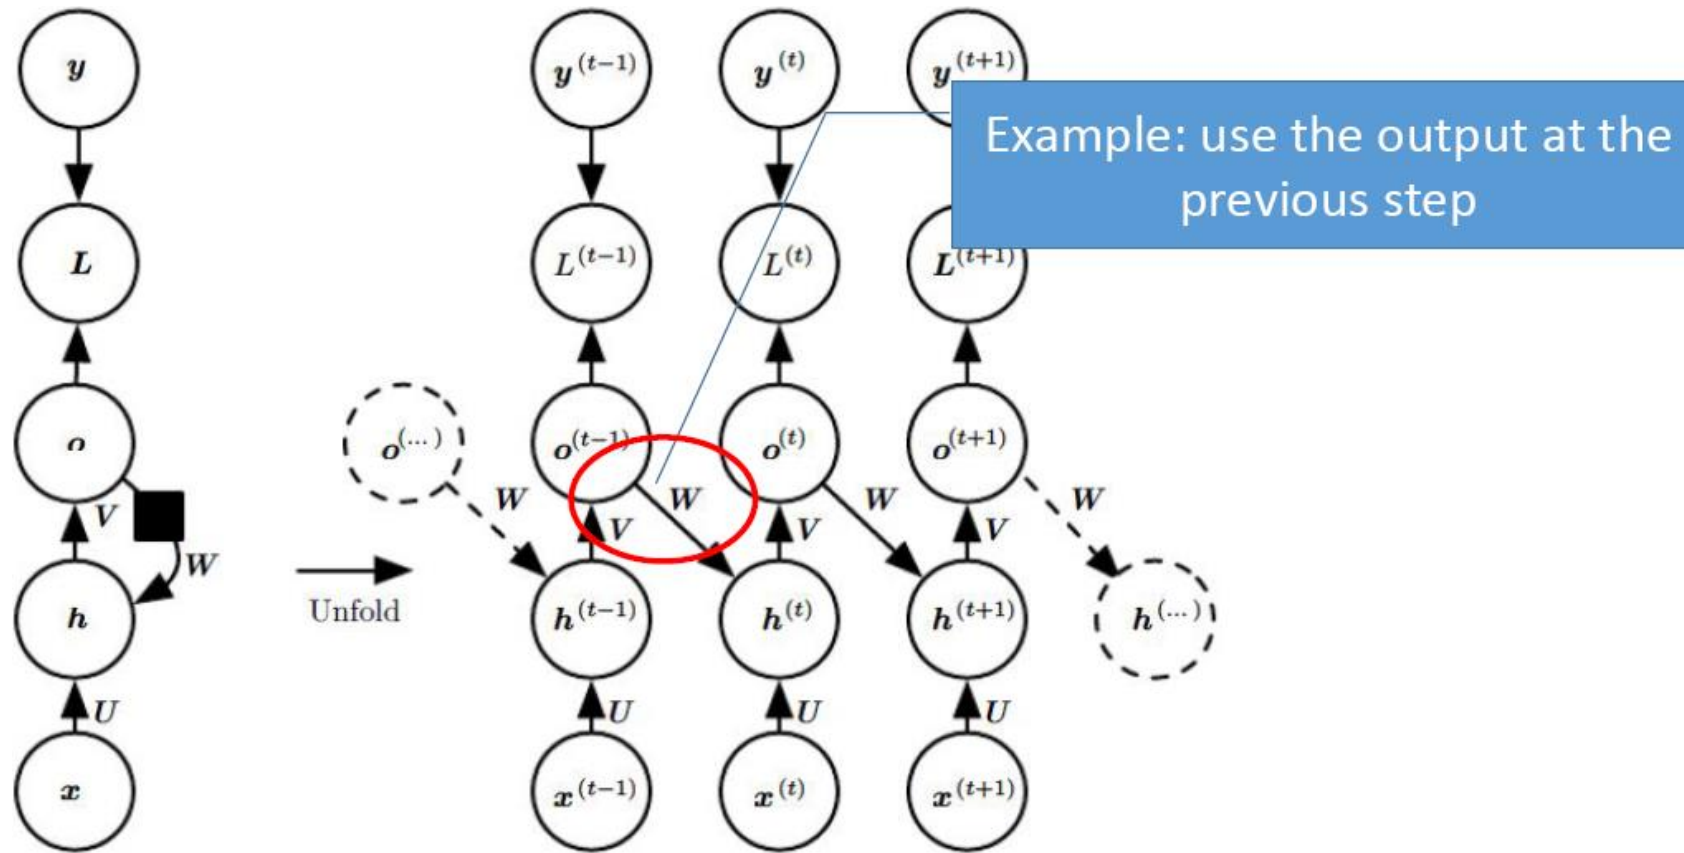


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

RNN Variations (2/2)

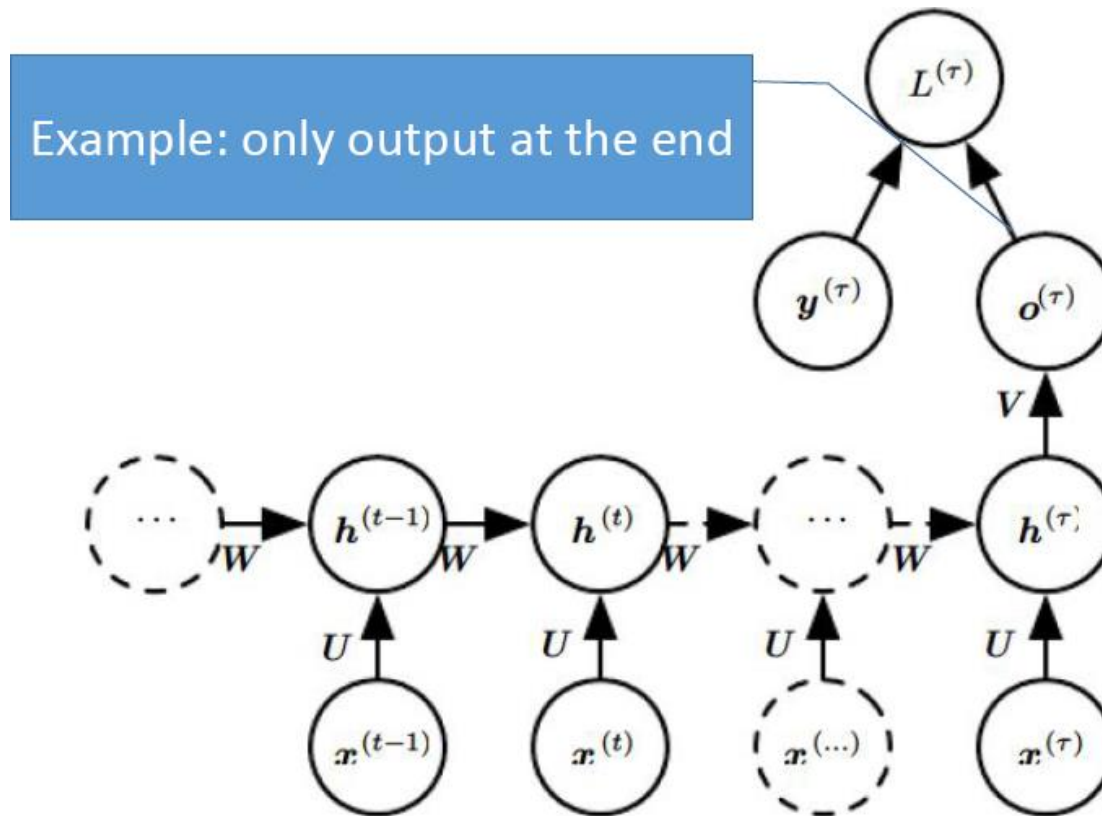


Figure from *Deep Learning*, Goodfellow, Bengio and Courville

Bidirectional RNNs (1/2)

- Many applications: output at time t may depend on the whole input sequence
- Example in speech recognition: correct interpretation of the current sound may depend on the next few phonemes, potentially even the next few words
- Bidirectional RNNs are introduced to address this

BiRNNs (2/2)





Image source: *Deep Learning*, Goodfellow, Bengio and Courville

Multilayer RNNs



- Multilayer RNNs

$$h_t^l = \tanh W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$

$$h \in \mathbb{R}^n, \quad W^l [n \times 2n]$$

- LSTM

$$W^l [4n \times 2n]$$

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \text{sigm} \\ \text{sigm} \\ \text{sigm} \\ \text{tanh} \end{pmatrix} W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$

$$c_t^l = f \odot c_{t-1}^l + i \odot g$$

$$h_t^l = o \odot \tanh(c_t^l)$$

Encoder-decoder RNNs (1/2)

- RNNs: can map sequence to one vector; or to sequence of same length
- What about mapping sequence to sequence of different length?
- Example: speech recognition, machine translation, question answering, etc.

Encoder-decoder RNNs (2/2)




Image source: *Deep Learning*, Goodfellow, Bengio and Courville